Course Reminders

Due Dates:

- Nothing due this week
- A4 due *next* Sunday (5/26)

Notes:

- Inference Case Study in Section this week
  - Intentionally unguided - use notes from glass to guide thought
  - Less-structured than assignments
- Guest Lecture Reminders:
  - Friday 5/24 (Dr. Gina Merchant)
  - Friday 5/31 (Dr. Brad Voytek)

Machine Learning: Basics

Shannon E. Ellis, Ph.D
UC San Diego

Department of Cognitive Science
sellis@ucsd.edu

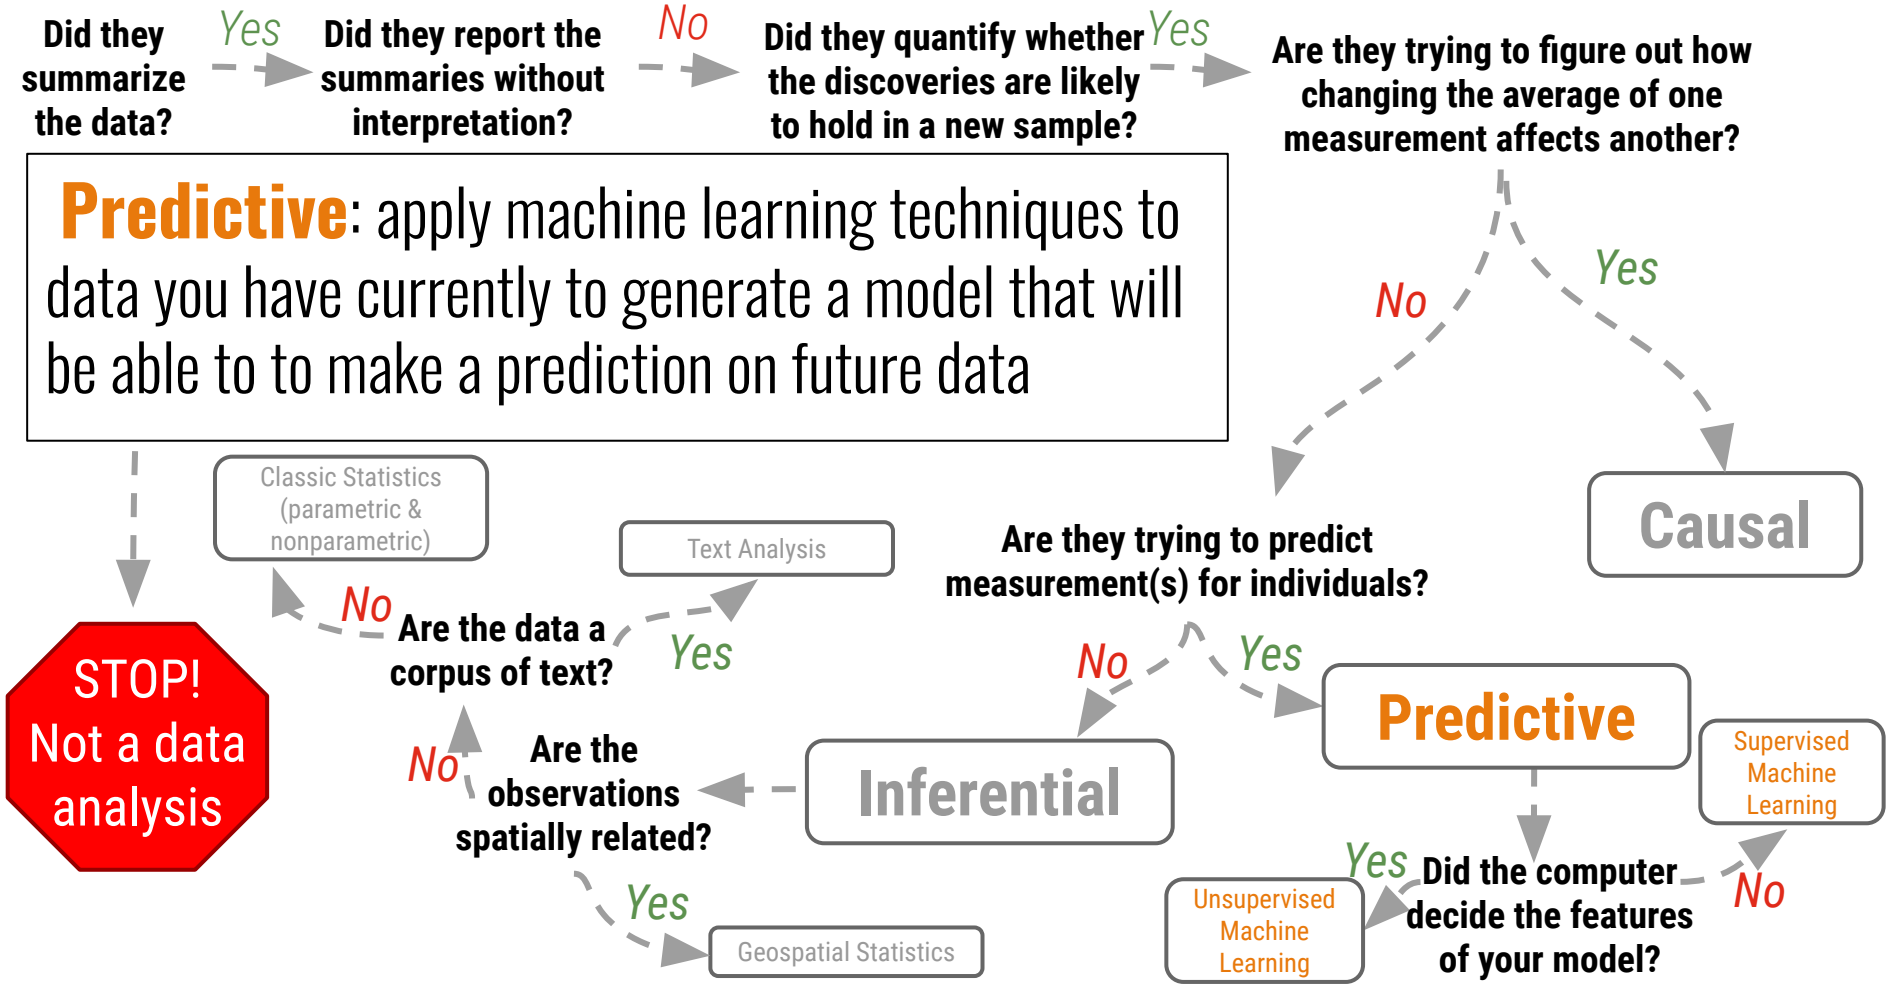
Did they summarize the data?
Yes
Did they report the summaries without interpretation?
No
Did they quantify whether the discoveries are likely to hold in a new sample?
Yes
Are they trying to figure out how changing the average of one measurement affects another?
Predictive: apply machine learning techniques to data you have currently to generate a model that will be able to to make a prediction on future data
No
Yes
Causal
Are they trying to predict measurement(s) for individuals?
No
Yes
Predictive
Supervised Machine Learning
Did the computer decide the features of your model?
Yes
Unsupervised Machine Learning
No
Inferential
Are the observations spatially related?
No
Yes
Geospatial Statistics
Are the data a corpus of text?
No
Yes
Text Analysis
Classic Statistics (parametric & nonparametric)
STOP! Not a data analysis

- **Problem:** Detecting whether credit card charges are fraudulent.
- **Data science question:** Can we use the time of the charge, the location of the charge, and the price of the charge to predict whether that charge is fraudulent or not?
- **Type of analysis:** Predictive analysis

predictive analysis uses data you have now to make predictions in the future

machine learning approaches are used for predictive analysis!

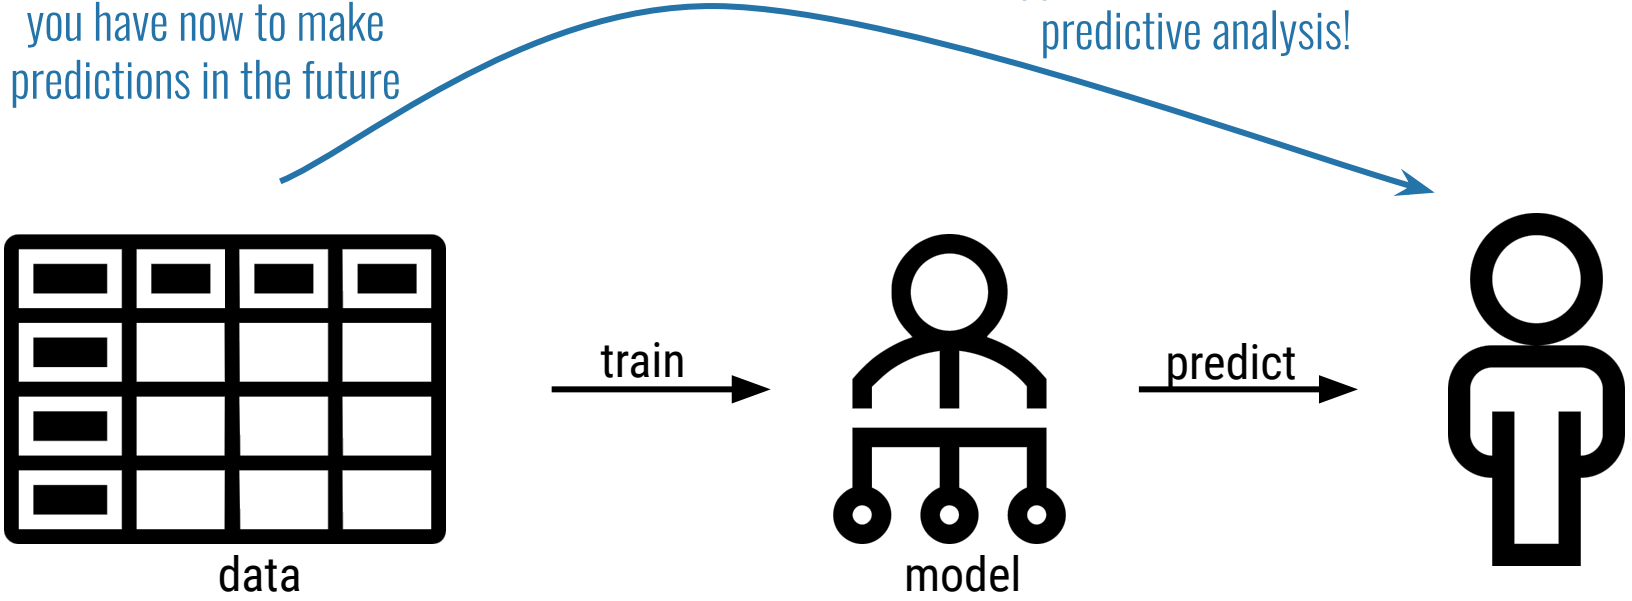

What is machine learning?

"Machine learning is the science of getting computers to act without being explicitly programmed"

- Andrew Ng, Stanford, ex-Google, chief scientist at Baidu, Coursera founder, Stanford Adjunct Faculty

Prediction Questions

Which of these questions is most appropriate for machine learning?

A How common is watching Sesame Street in the US?

B What is the effect of watching Sesame Street on children's brains?

C What is the relationship between early childhood educational programming and success in elementary school?

D Can we use information about one's early childhood to predict their success in elementary school?

E How does Sesame Street cause an increase in educational attainment?

Machine Learning Generalizations

Basic Steps to Prediction

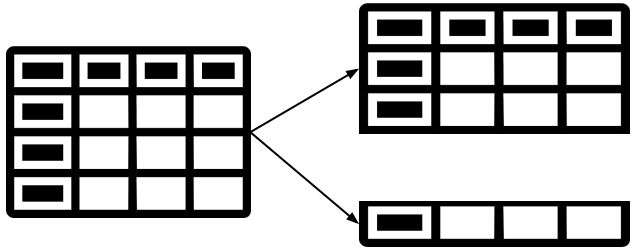

data
partitioning

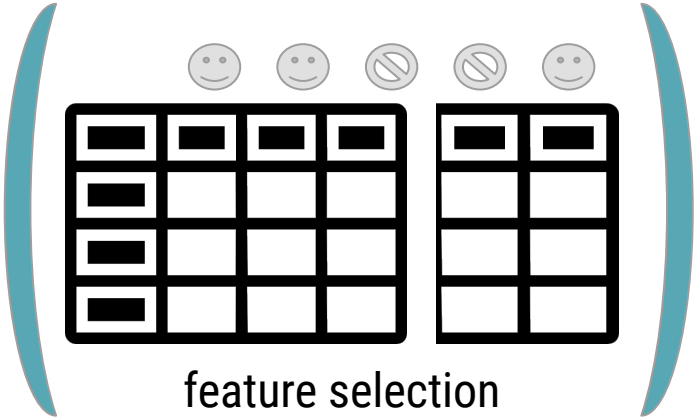

feature selection

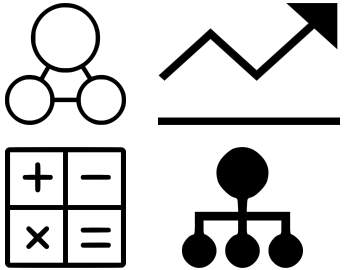

model selection

model assessment

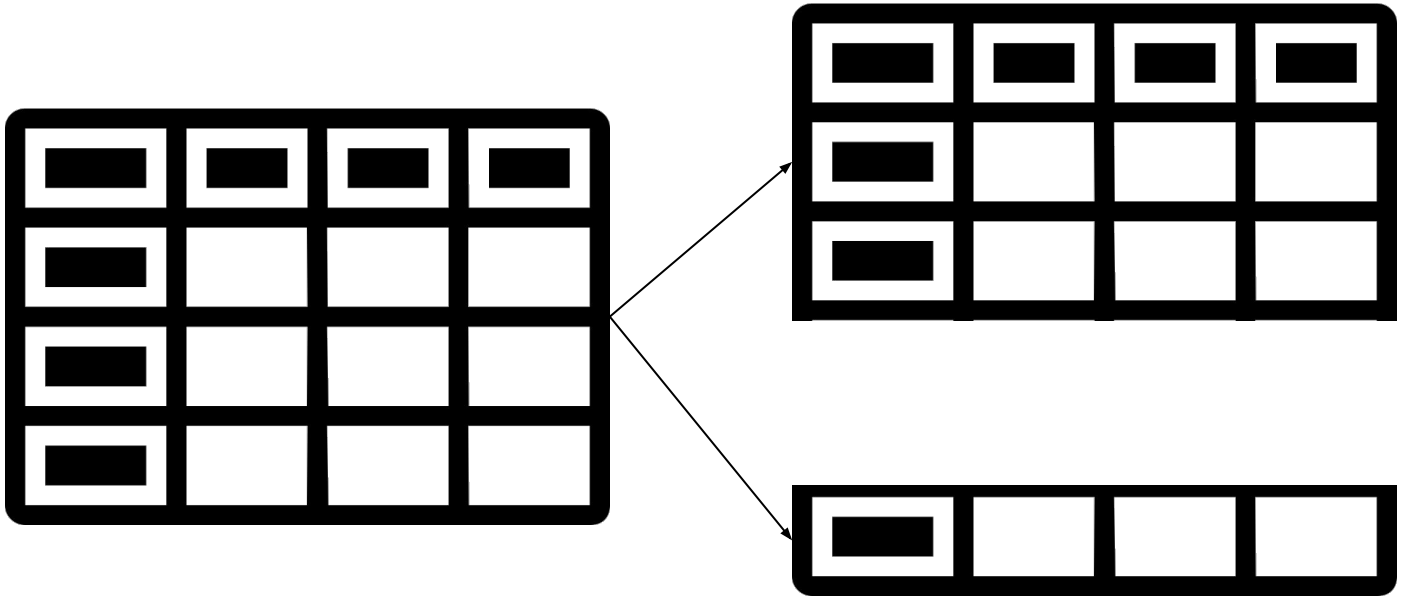

data partitioning

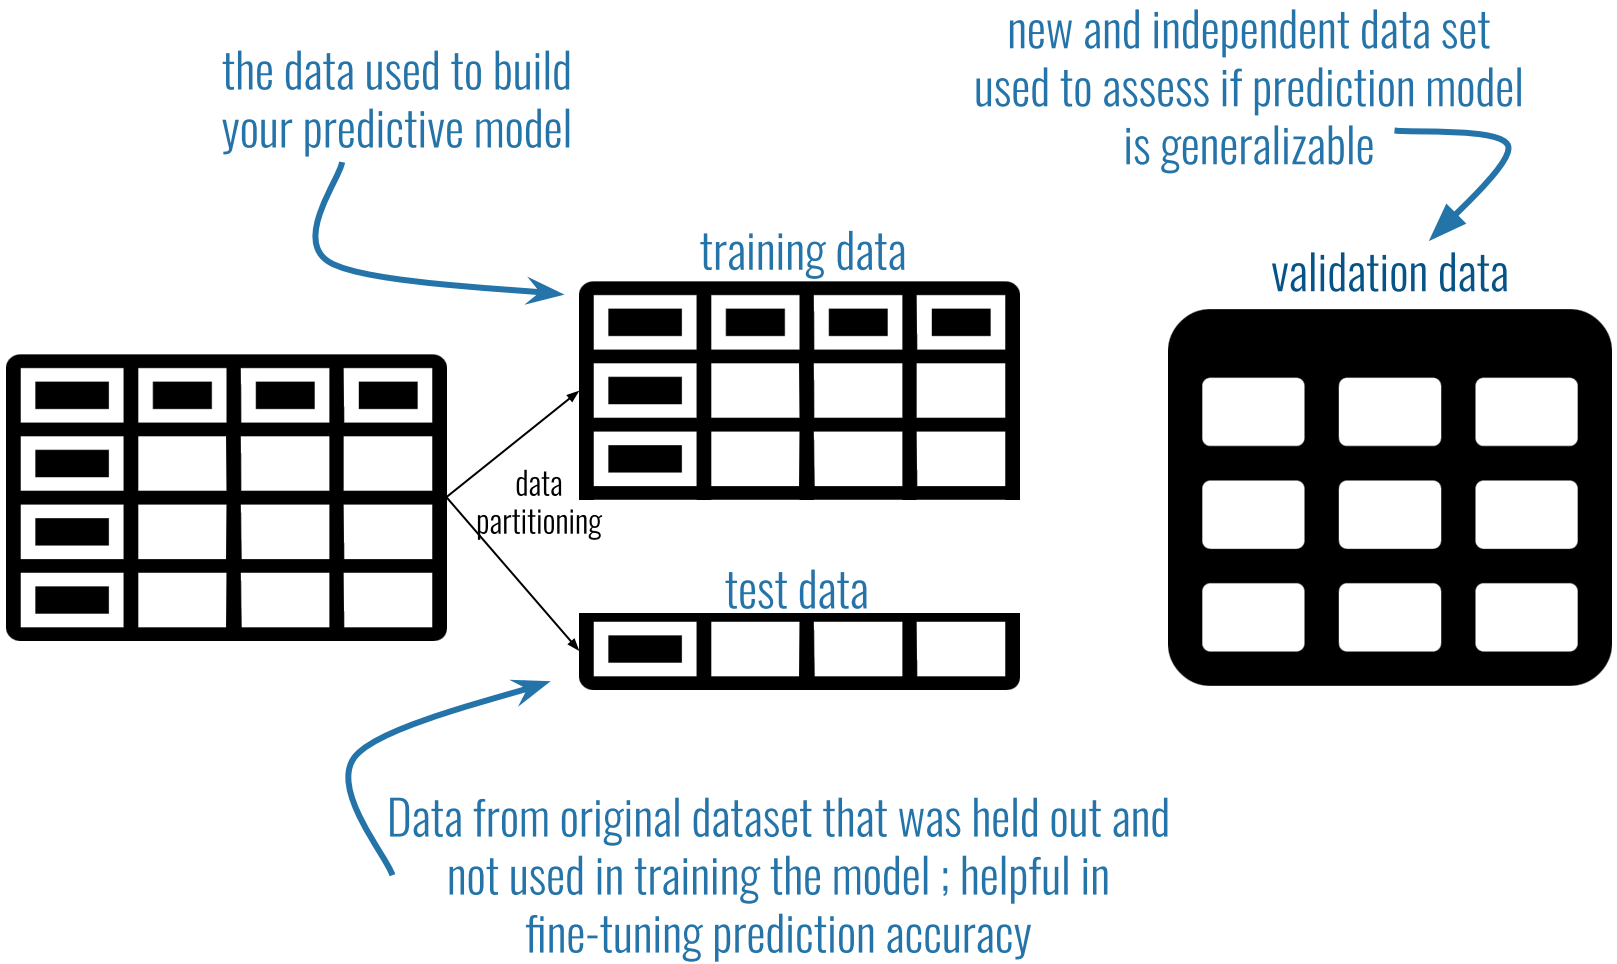
the data used to build your predictive model
new and independent data set used to assess if prediction model is generalizable
training data
validation data
data partitioning
test data
Data from original dataset that was held out and not used in training the model ; helpful in fine-tuning prediction accuracy

Data Partitioning

What portion of the data are typically used for generating the model?

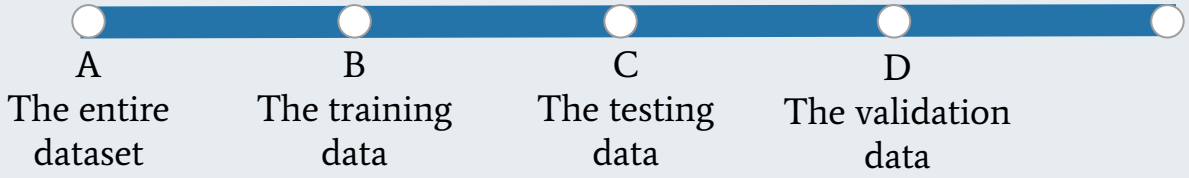

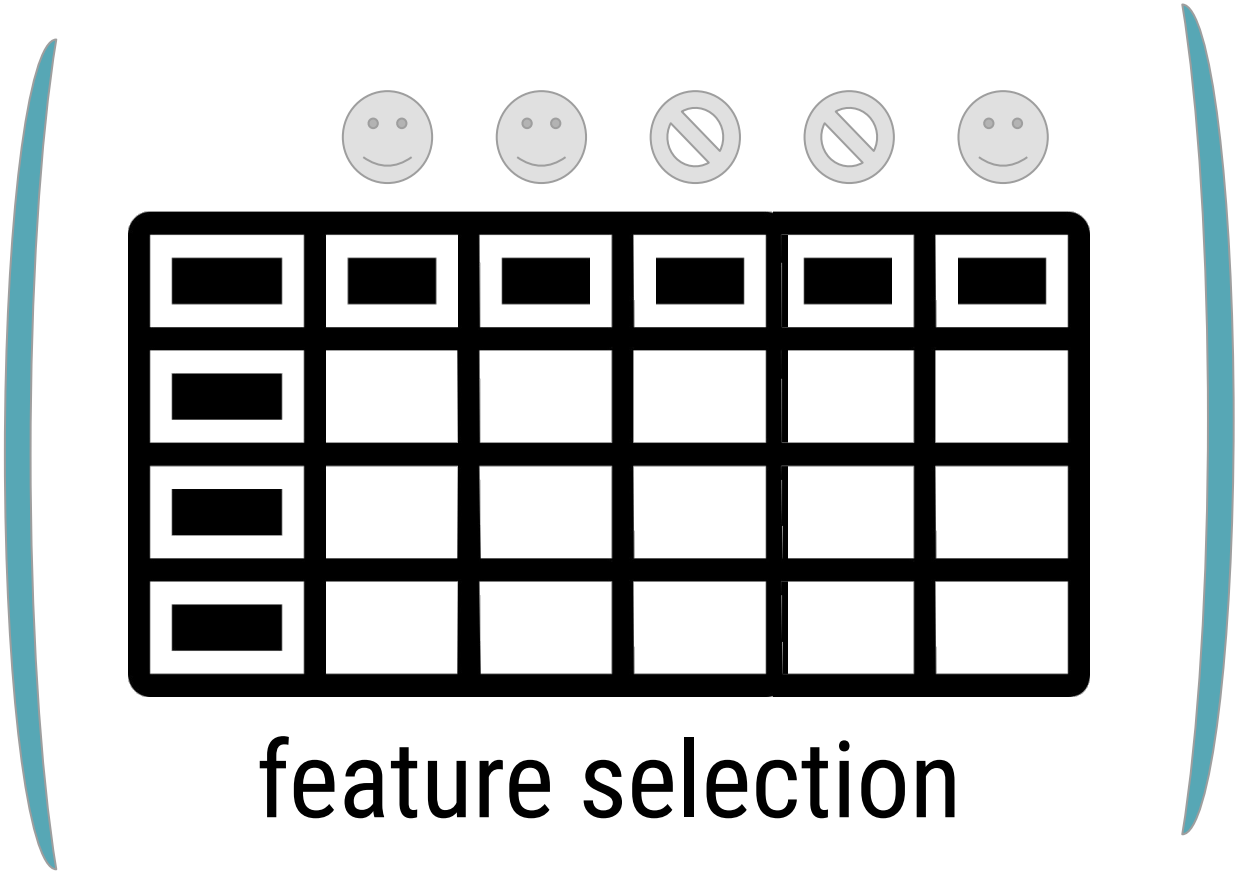
feature selection

elephant height data are likely
not predictive of US elections

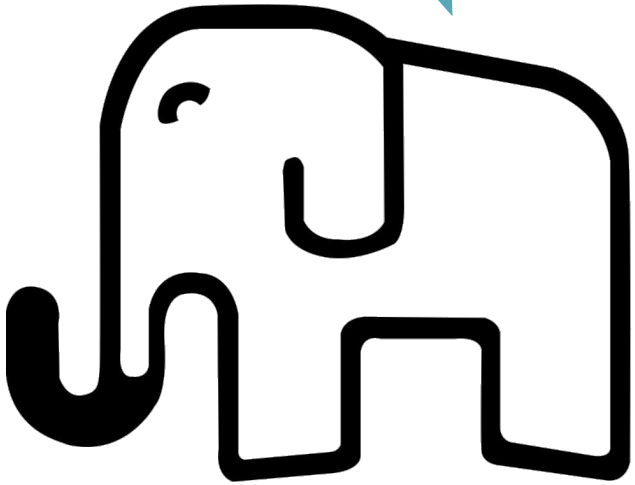

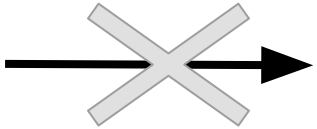

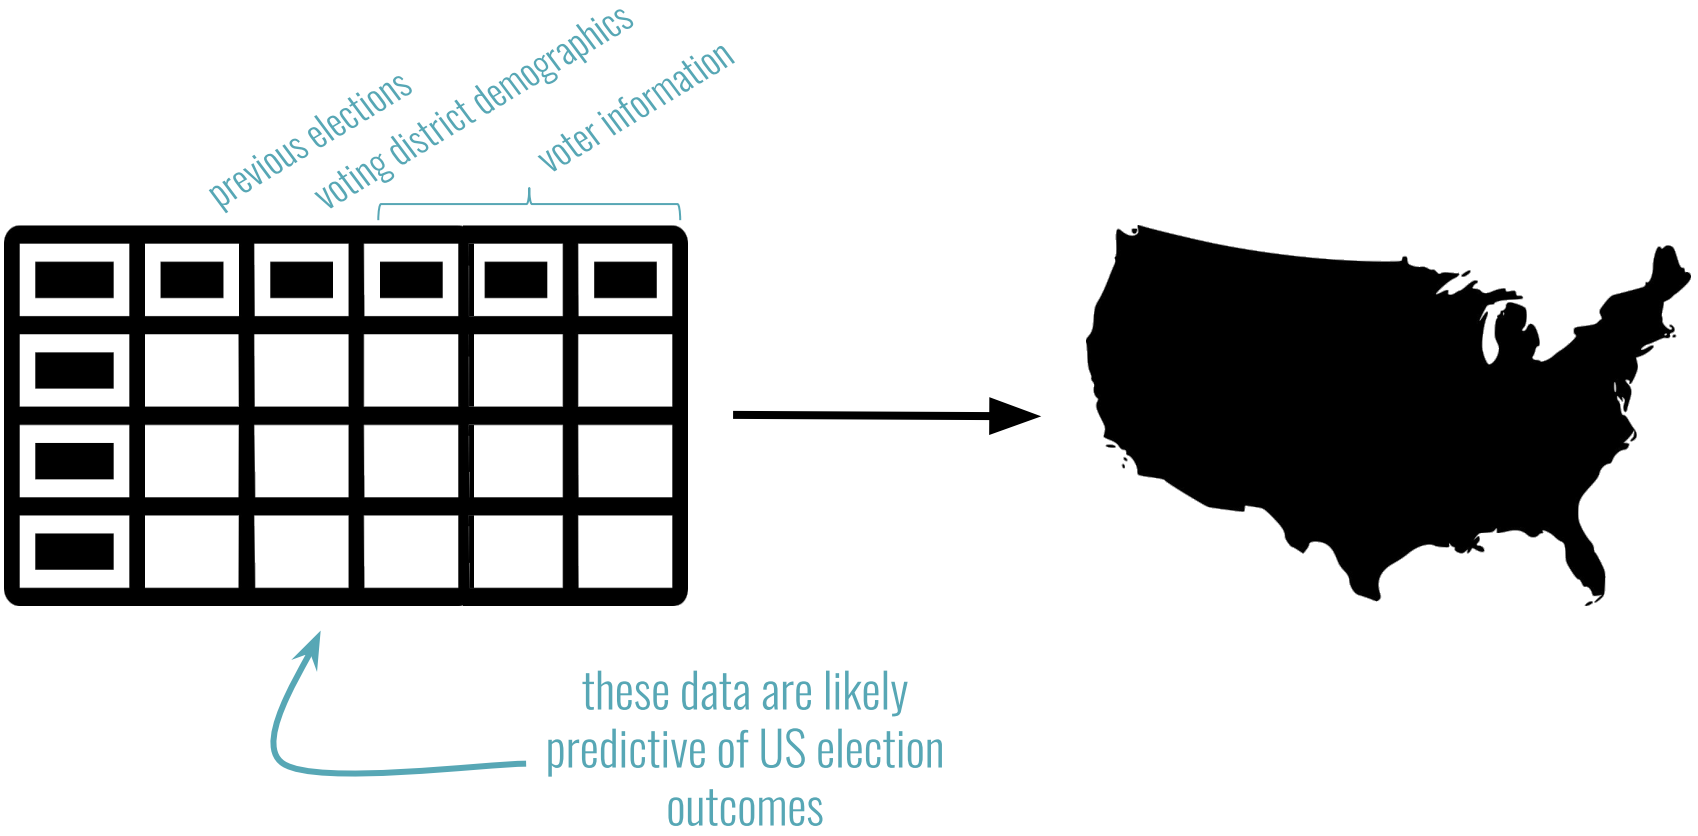
previous elections
voting district demographics
voter information
these data are likely predictive of US election outcomes

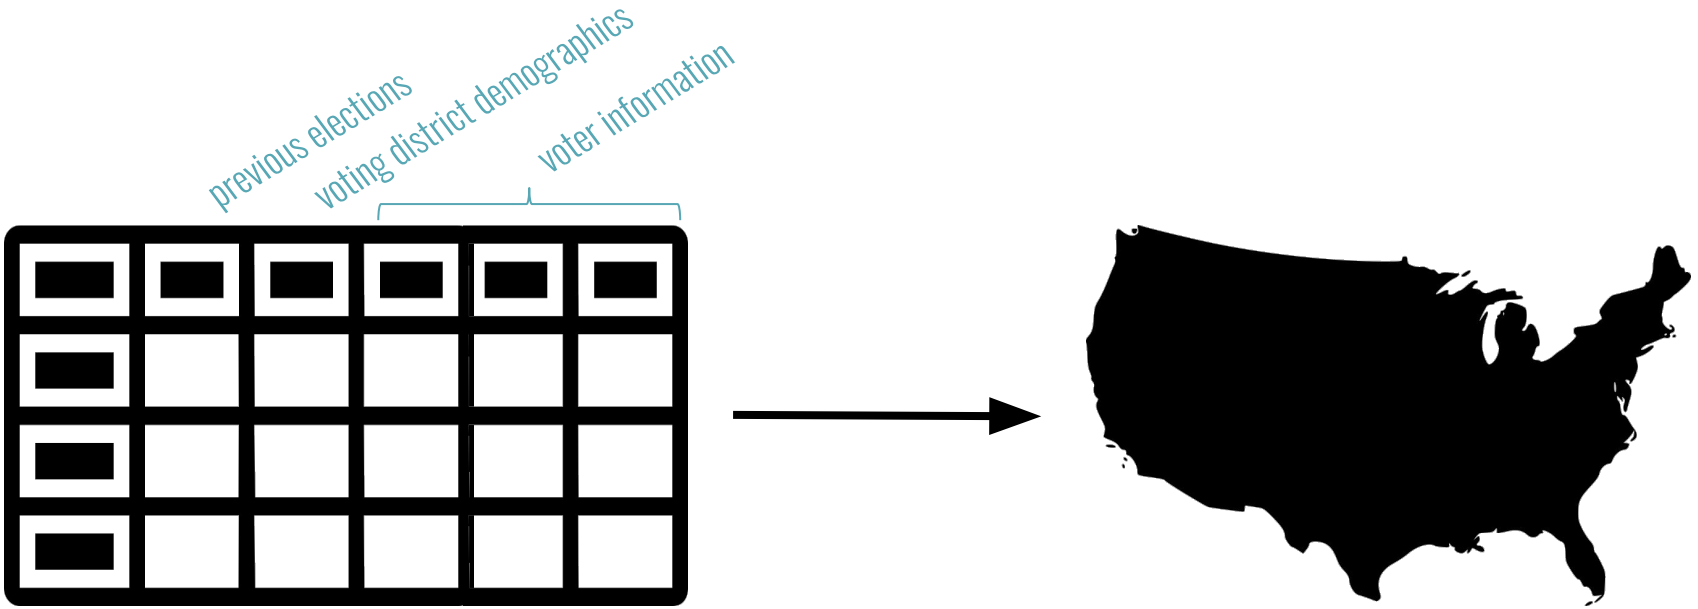

feature selection determines which variables are most predictive and includes them in the model

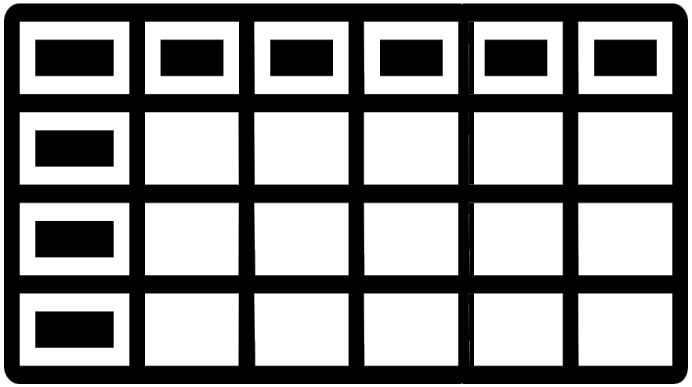

variables that can be used for accurate prediction exploit the relationship between the variables but do NOT mean that one causes the other

To modes of machine learning

Supervised Learning

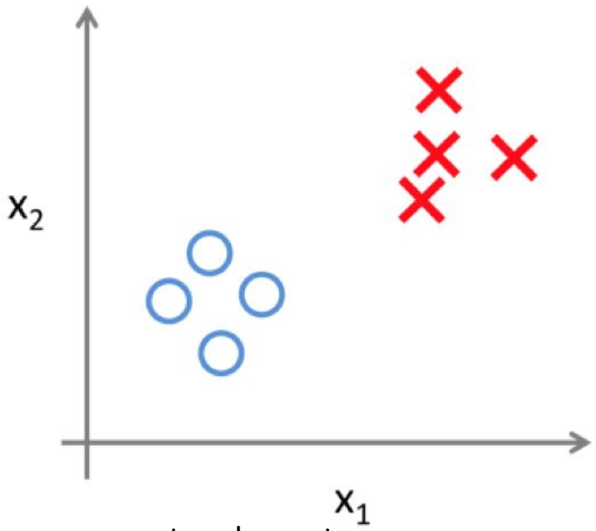

You tell the computer how to classify the observations

Unsupervised Learning

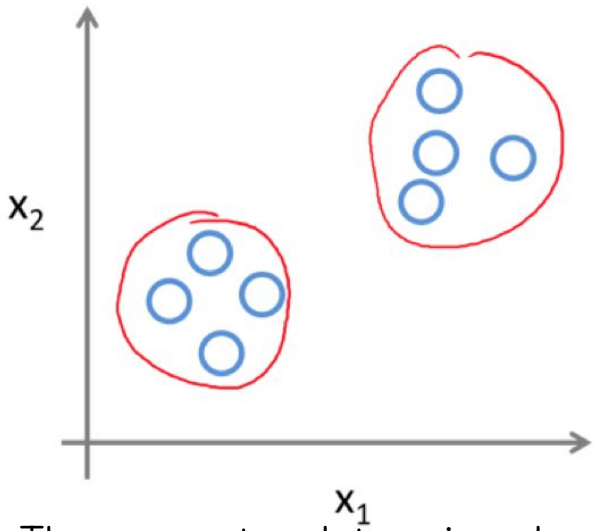

The computer determines how to classify based on properties within the data

Approaches to machine learning

Supervised Learning

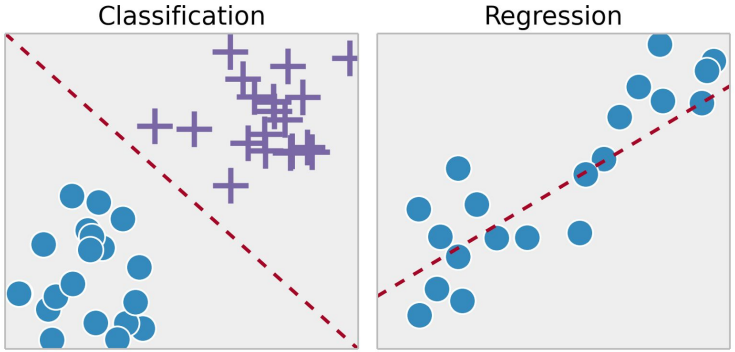

categorical variables

continuous variables

Prediction accuracy dependent on training data

Unsupervised Learning

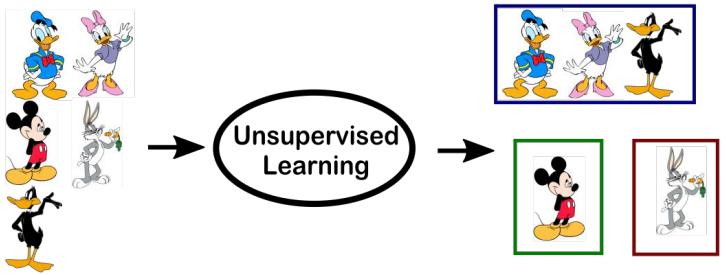

Clustering (categorical)
& dimensionality reduction (continuous)

can automatically identify structure in data

Image source: https://towardsdatascience.com/supervised-vs-unsupervised-learning-14f68e32ea8d

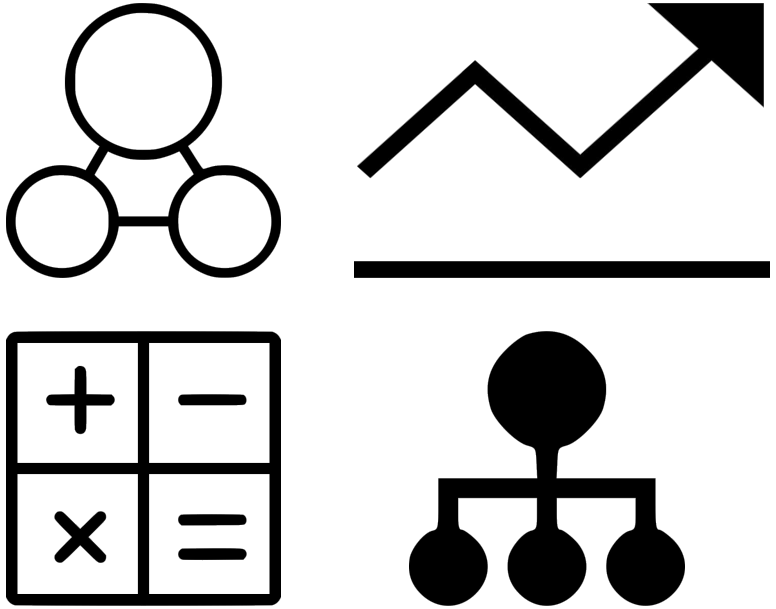

model selection

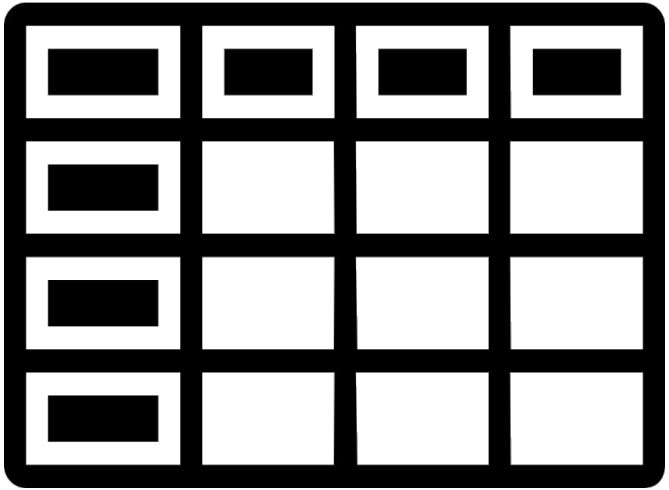

big
datasets

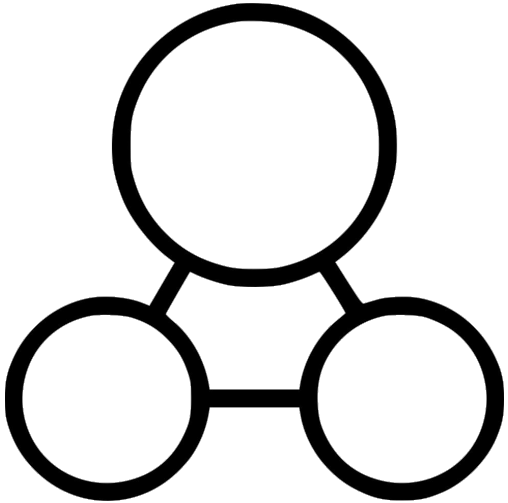

simple
models

Supervised Learning

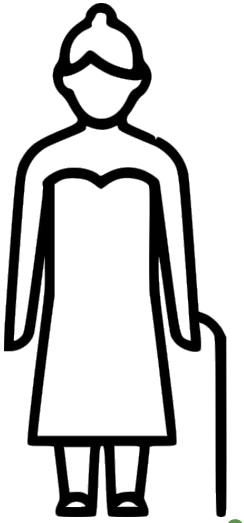

Regression:

predicting continuous variables
(i.e. Age)

continuous variable prediction

Classification:

predicting categorical variables
(i.e. education level)

categorical variable prediction

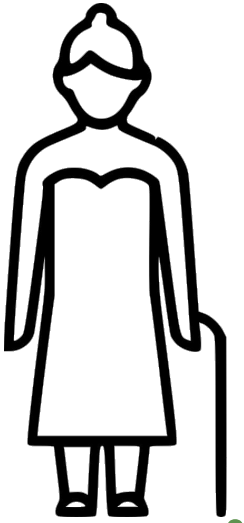

Regression:

predicting continuous variables
(i.e. Age)

continuous variable prediction

Classification:

predicting categorical variables
(i.e. education level)

categorical variable prediction

Supervised Learning

regression

continuous variable prediction

Supervised Learning

regression

continuous variable prediction

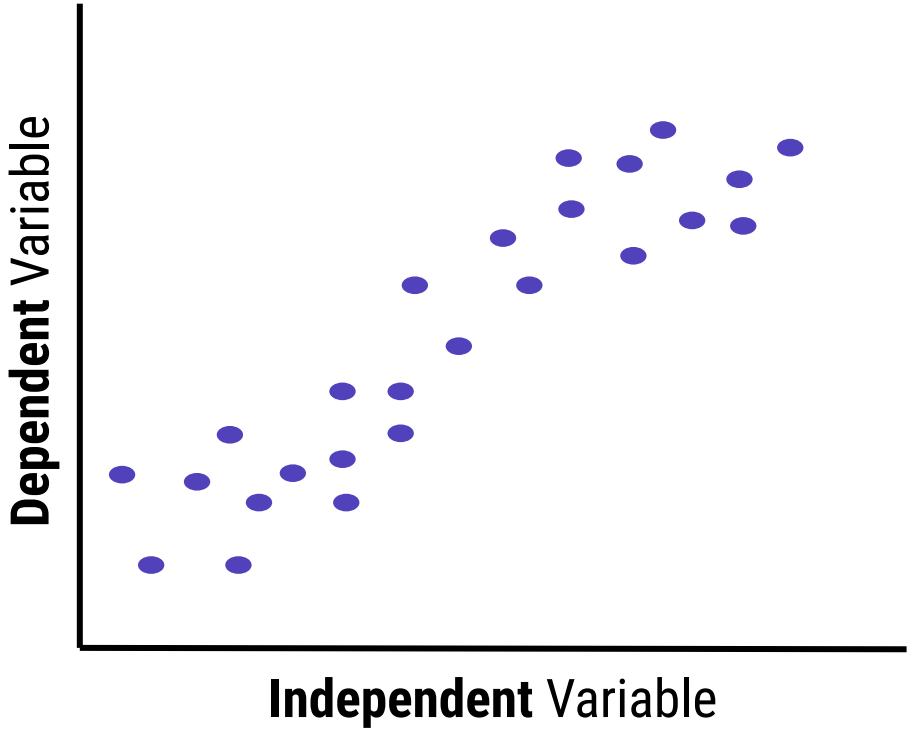

Supervised Learning

regression

continuous variable prediction

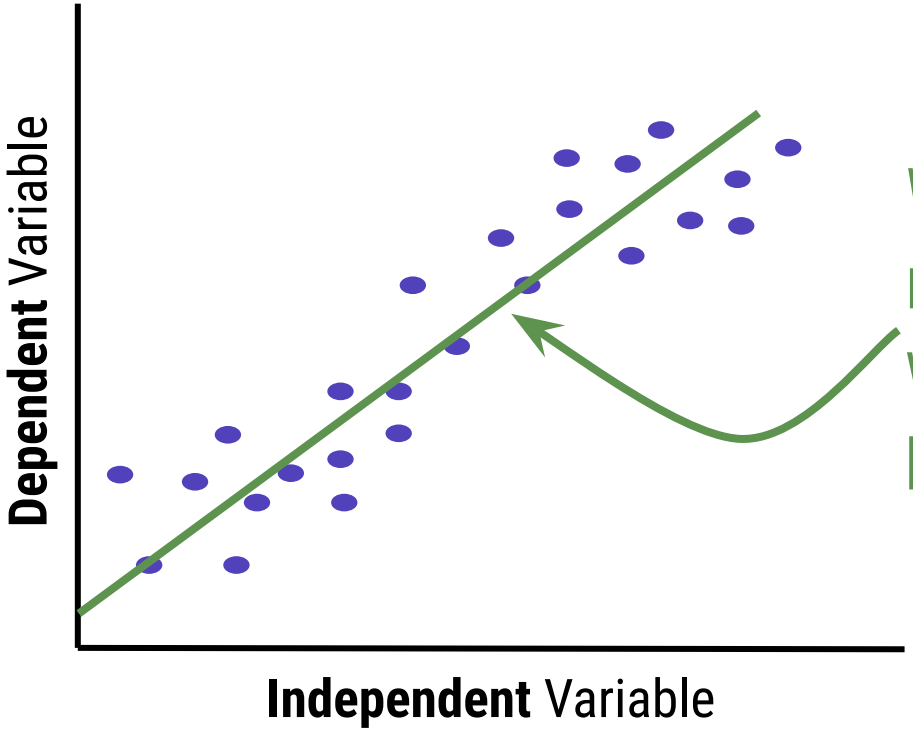

We'll use the linear relationship between variables to generate a **predictive model**

Supervised Learning

regression

continuous variable prediction

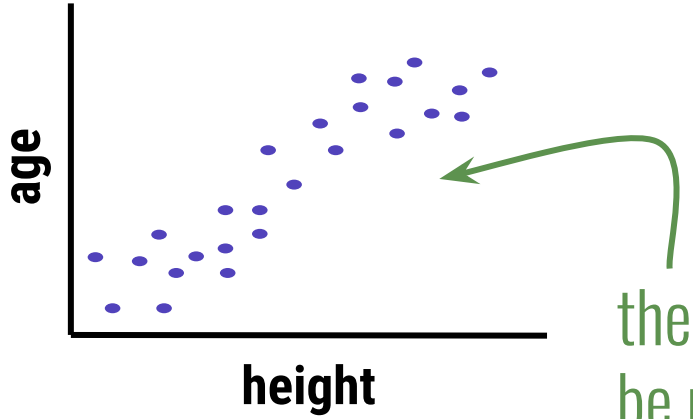

the training data will be used to build the predictive model

Supervised Learning

regression

continuous variable prediction

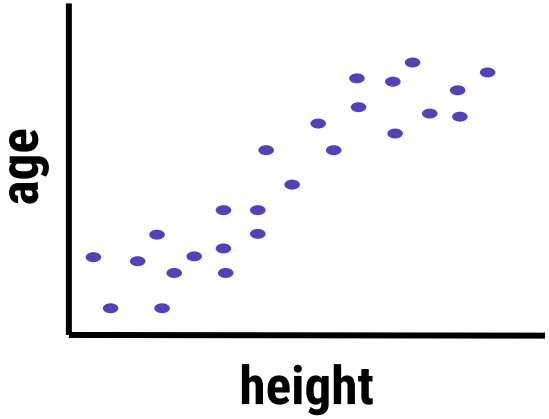

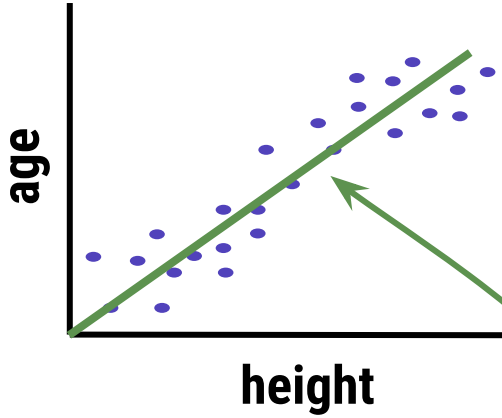

use linear regression to model the relationship

Supervised Learning

regression

continuous variable prediction

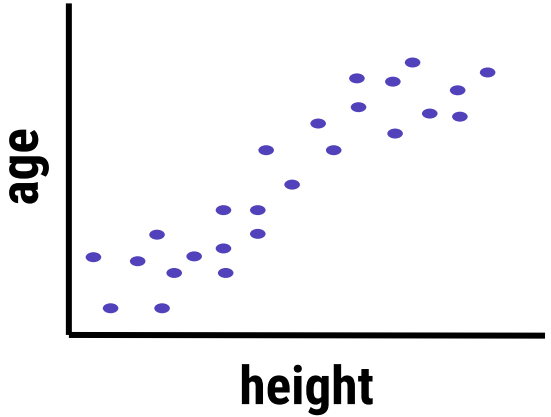

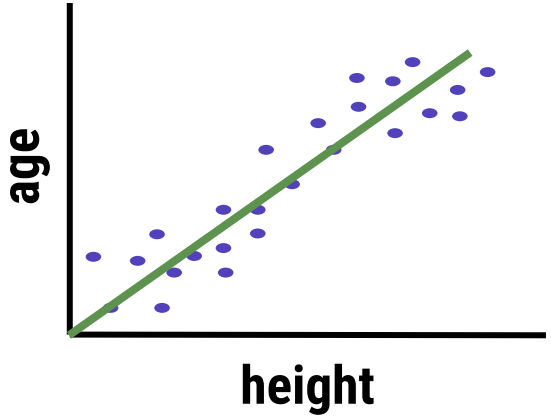

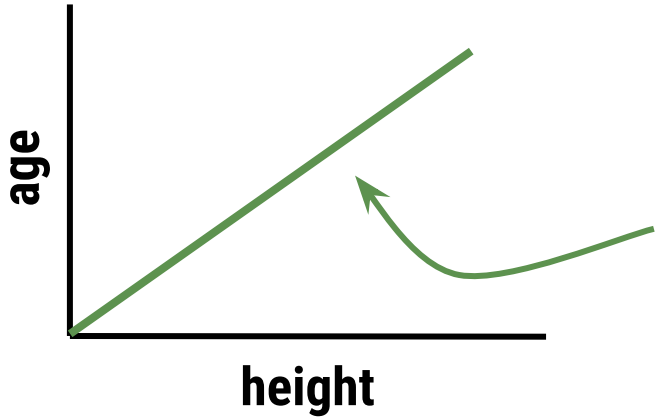

For prediction, the individual values in the training data are *not* important. We only need the model.

Supervised Learning

regression

continuous variable prediction

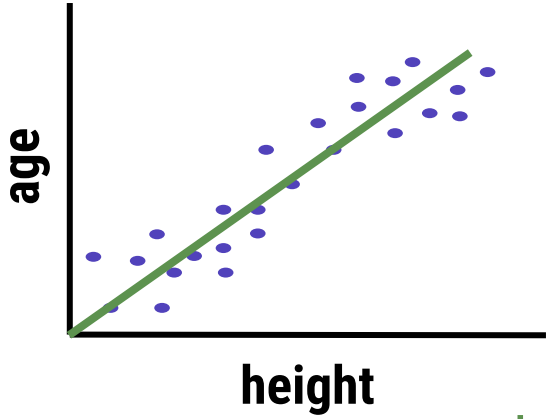

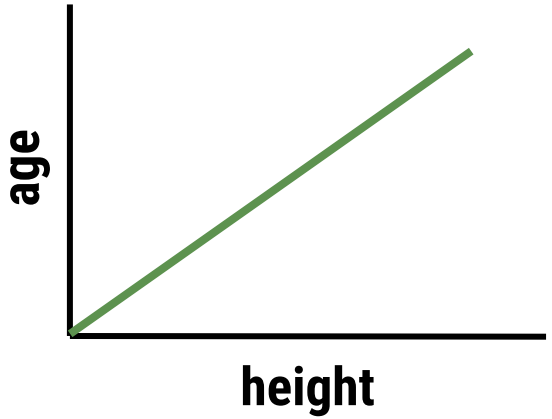

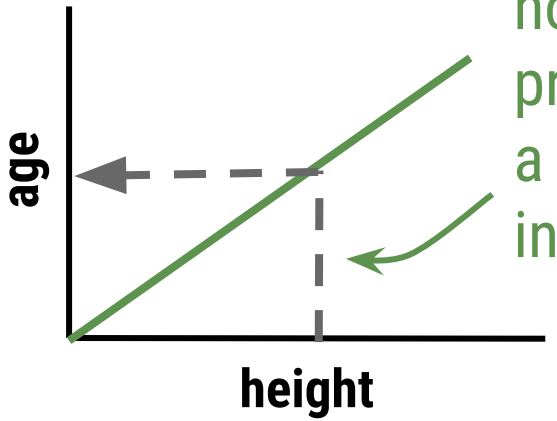

how we'll make predictions for a future individual

Supervised Learning

regression

continuous variable prediction

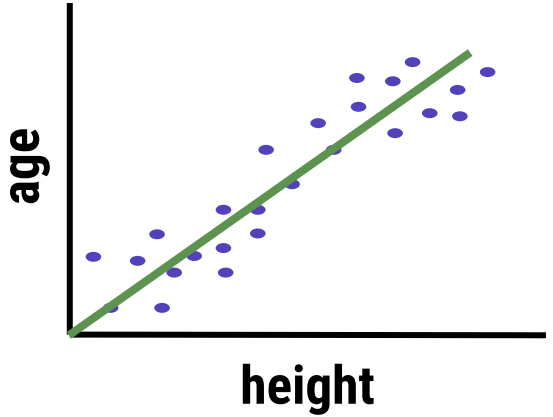

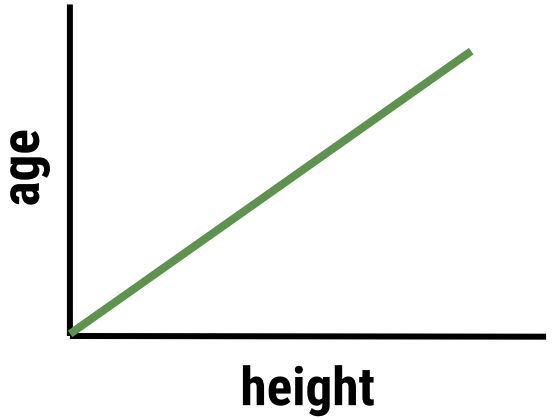

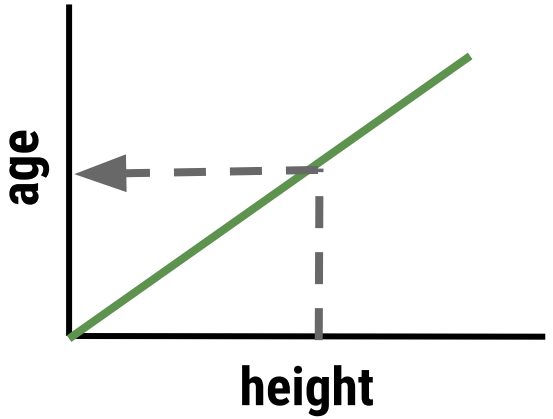

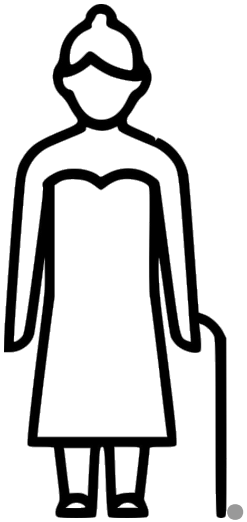

Regression:

predicting continuous variables
(i.e. Age)

Classification:

predicting categorical variables
(i.e. education level)

Supervised Learning

classification

categorical variable prediction

Supervised Learning

classification

categorical variable prediction

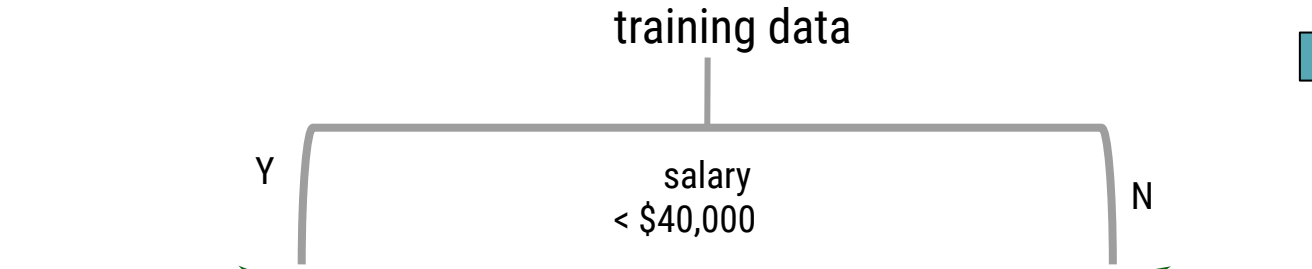

All the people who make *less* than 40K over here

All the people who make *more* than 40K over here

Supervised Learning

classification

categorical variable prediction

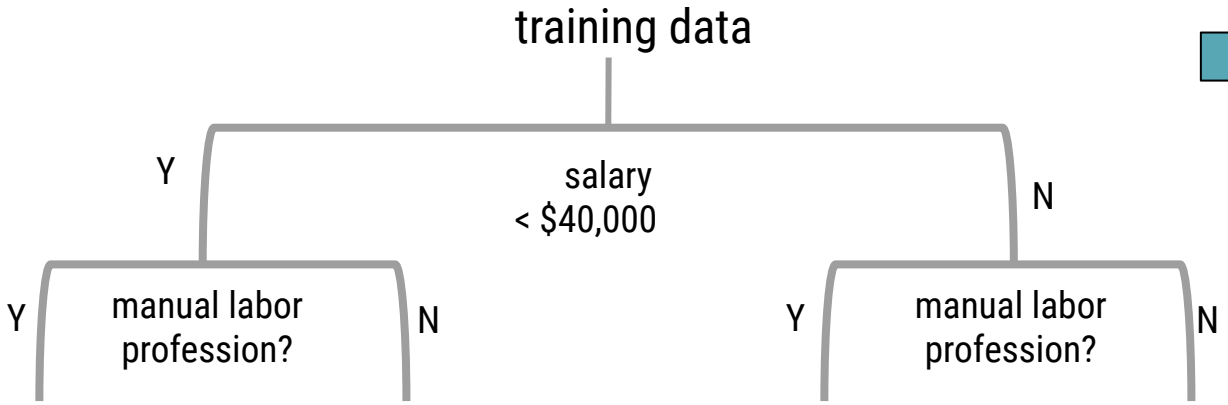

Continue adding *branches* to the **decision tree** where the variables and information in the training data decide which observations goes down which branch

Supervised Learning

classification

categorical variable prediction

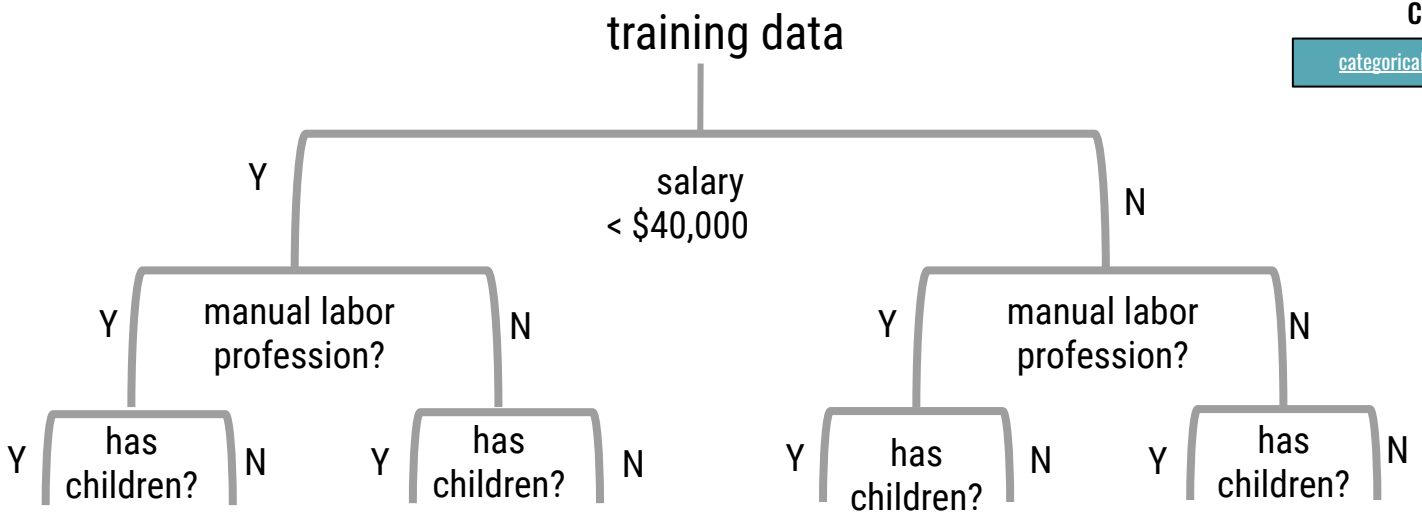

Supervised Learning

classification

categorical variable prediction

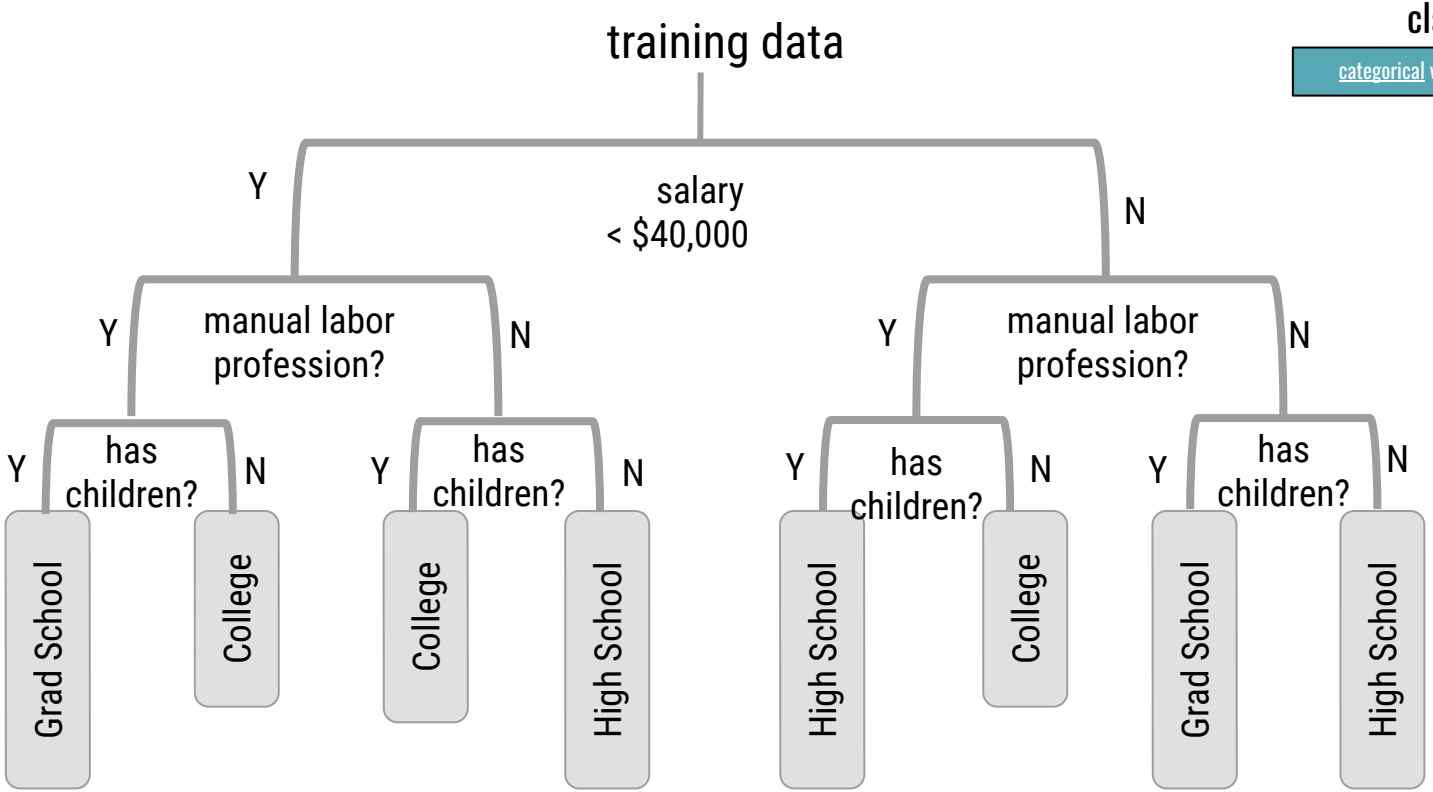

At the end of the tree, labels will be applied to each *leaf* of the tree

Supervised Learning

classification

categorical variable prediction

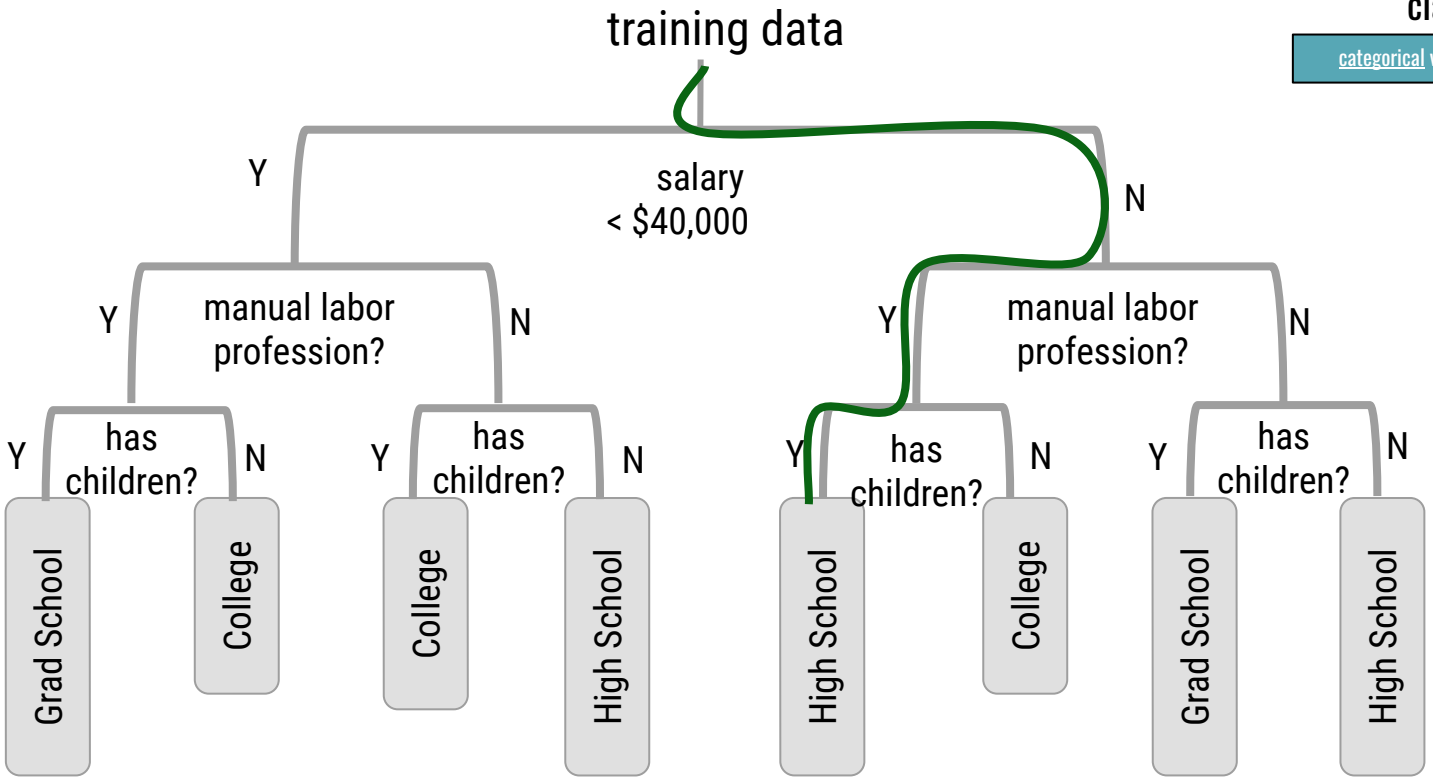

Decision trees are helpful for classification

Supervised Learning

classification

categorical variable prediction

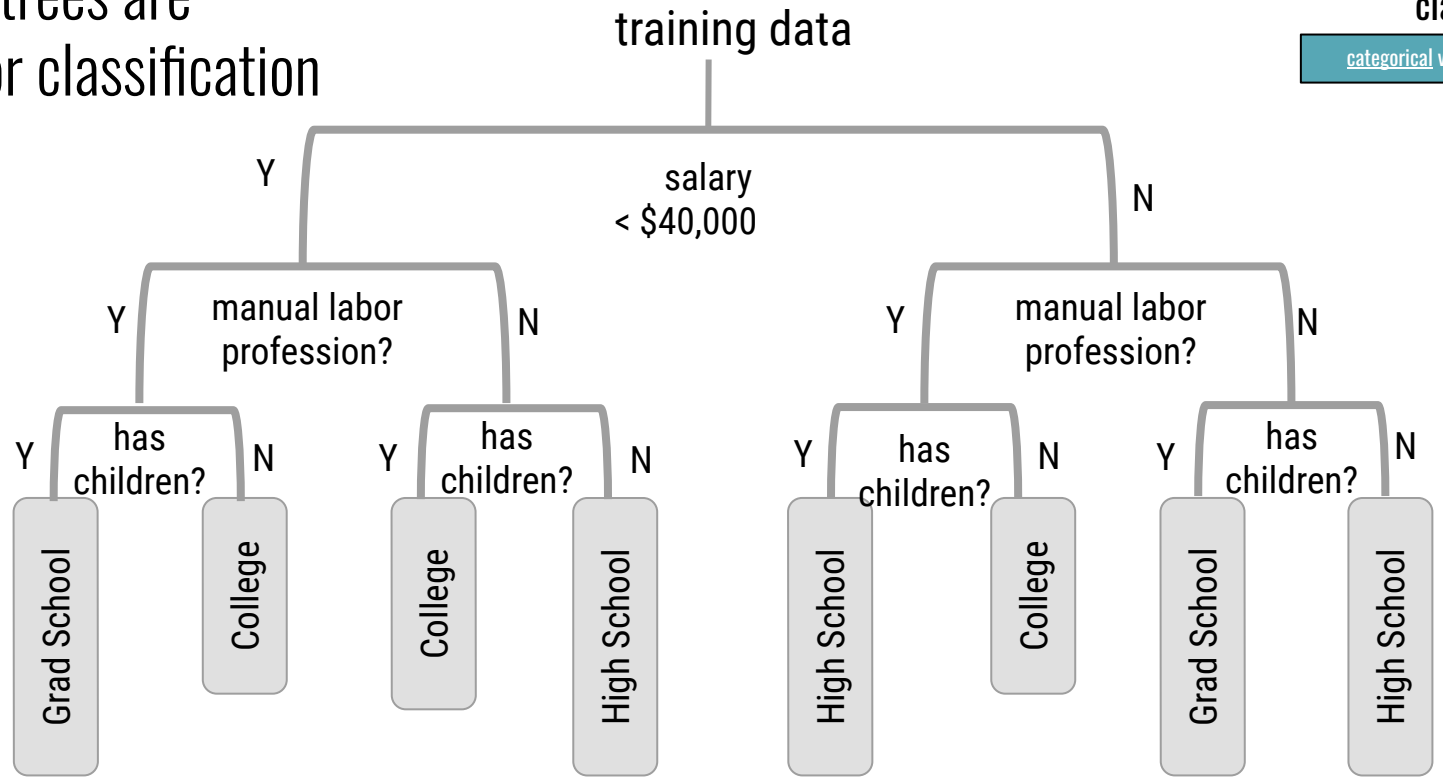

Unsupervised Learning

To modes of machine learning

Supervised Learning

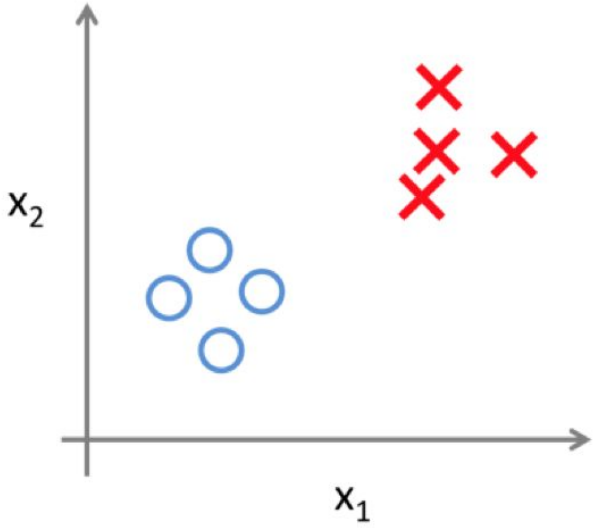

Unsupervised Learning

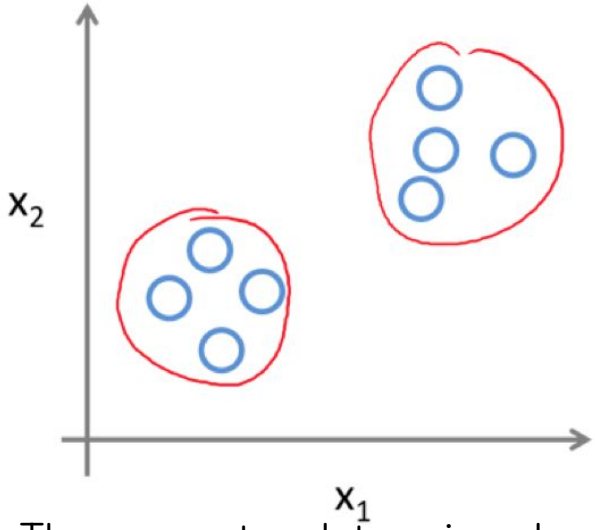

The computer determines how to classify based on properties within the data

Features are given as inputs

model identifies patterns in the input data

predictions are output

No labels are provided during modeling

PCA, k-means clustering, t-SNE, neural nets, self-organizing maps, (i.e. facial recognition, image processing, EDA)

model identifies patterns
in the input data

Is it cake?

https://aczane.netlify.com/2018/02/08/the-first-and-namesake-post-is-it-cake/

model identifies patterns
in the input data

Is it cake?

1

2

https://aczane.netlify.com/2018/02/08/the-first-and-namesake-post-is-it-cake/

model identifies patterns in the input data

Is it cake?

1

2

new emoji

Is it cake?

No → Class 1 NOT CAKE

Yes → Class 2 CAKE

https://aczane.netlify.com/2018/02/08/the-first-and-namesake-post-is-it-cake/

WHAT IS A NEURON?

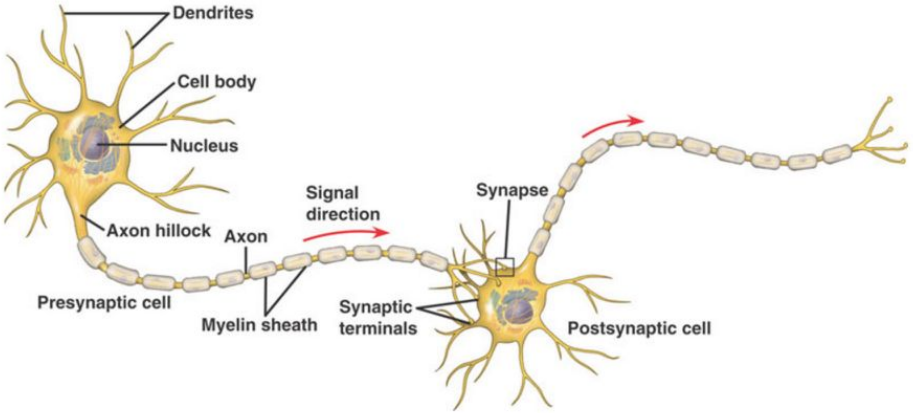

- Receives signal on synapse
- When trigger sends signal on axon

MATHEMATICAL NEURON

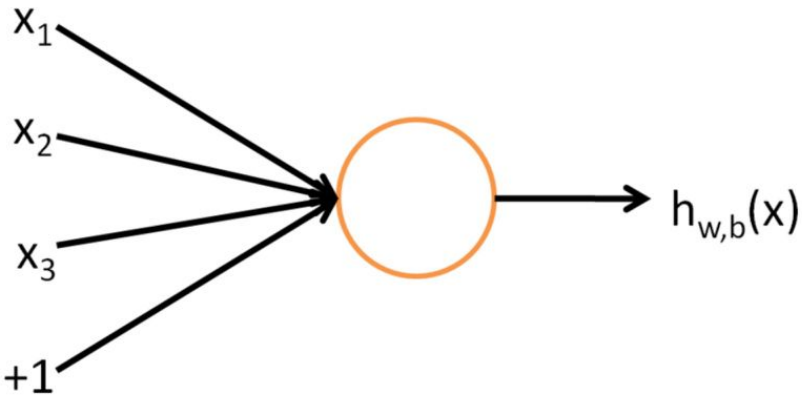

- Mathematical abstraction, inspired by biological neuron
- Either on or off based on sum of input

This will likely not be the last time you see this (mostly unhelpful) neural net image

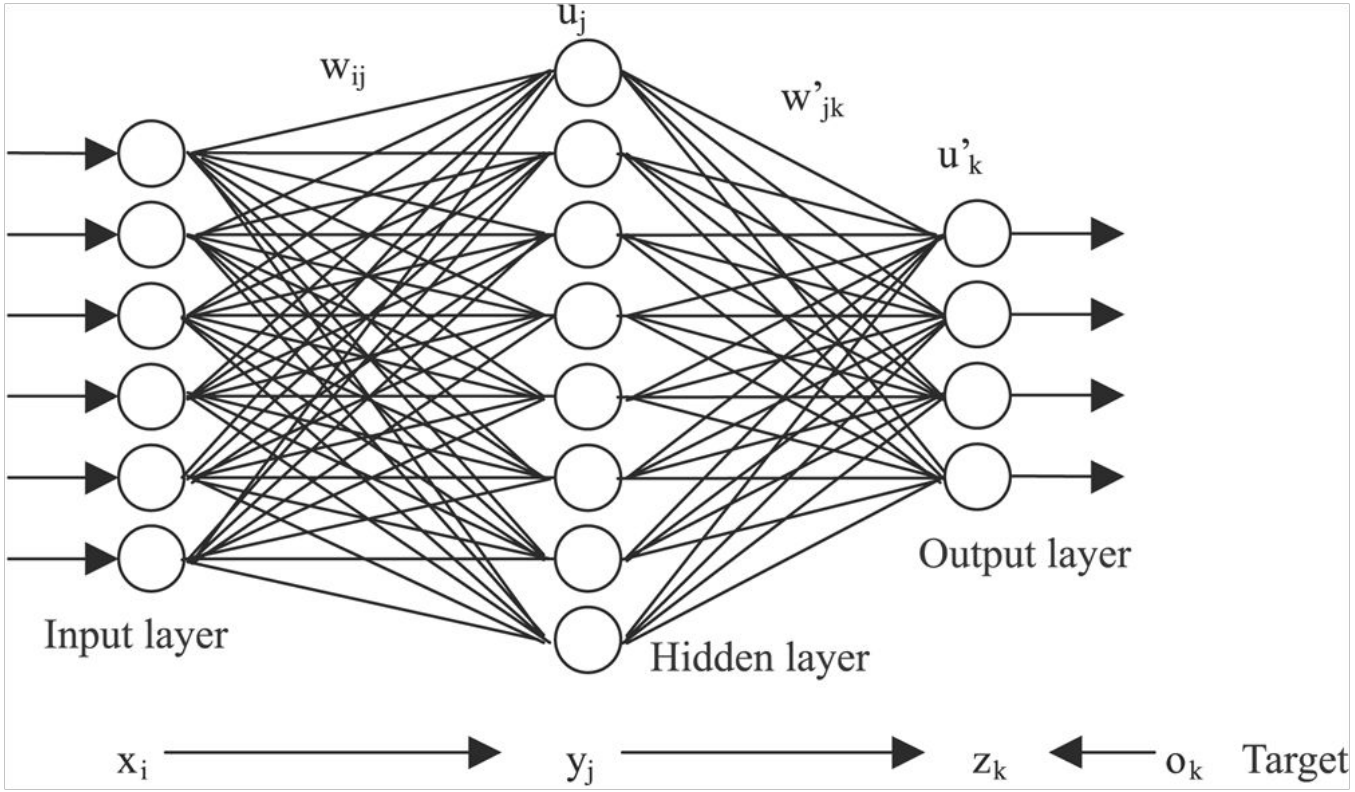

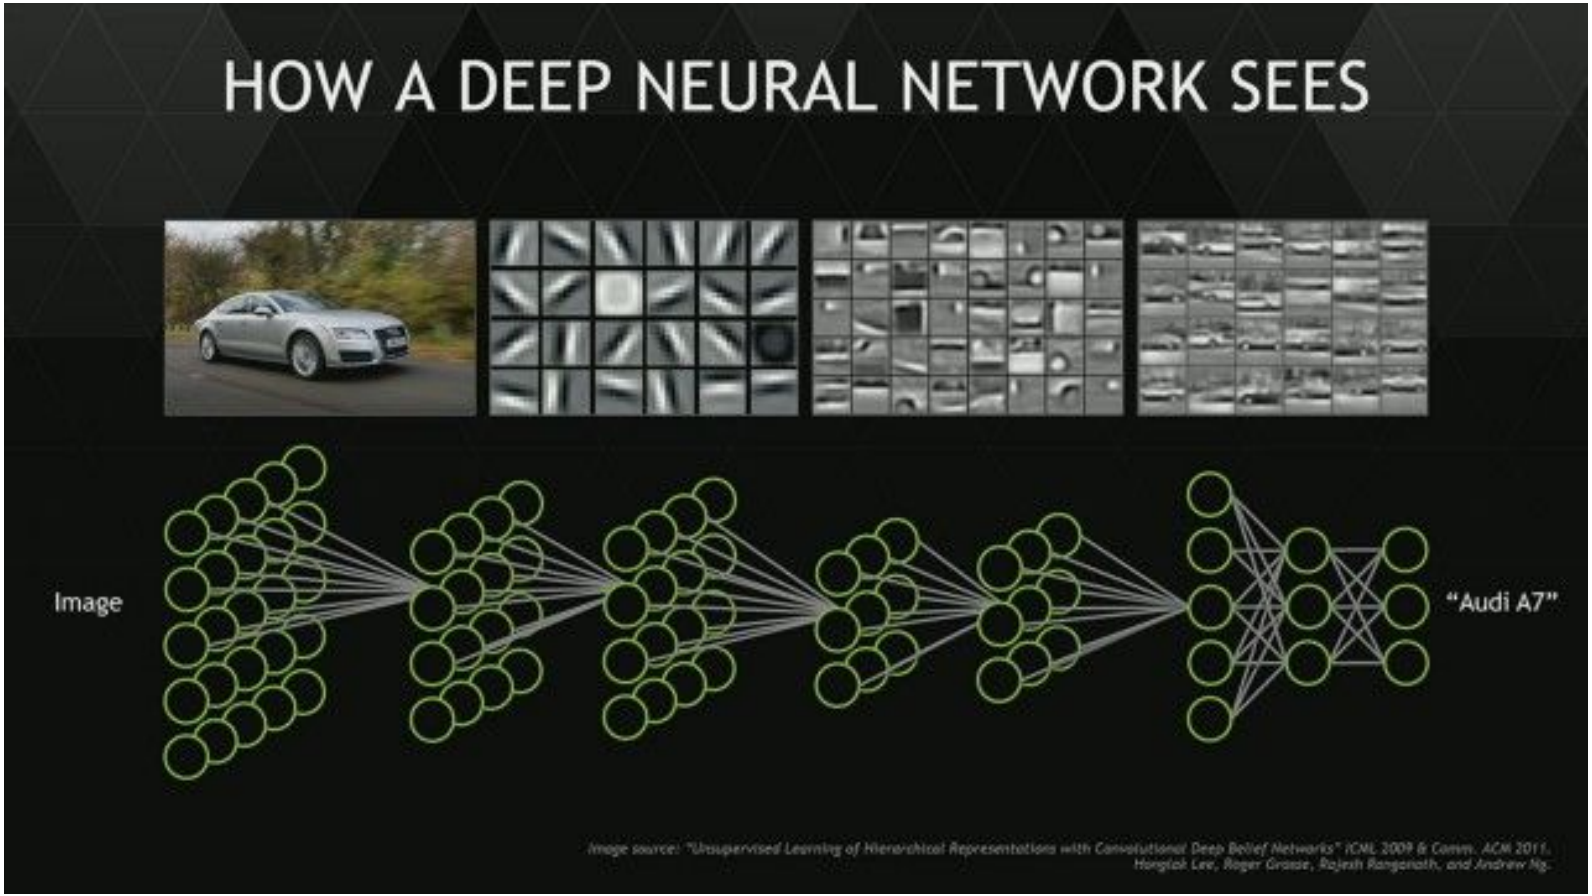

Image Source: https://towardsdatascience.com/understanding-residual-networks-9add4b664b03

Supervised Learning

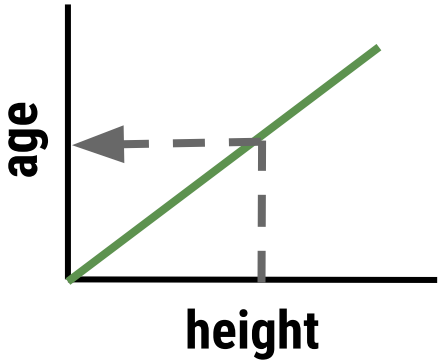

continuous

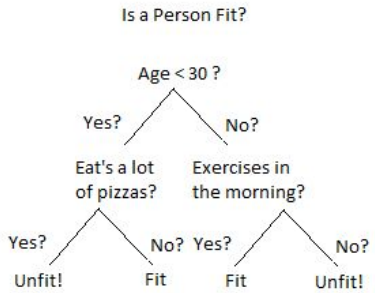

classification

Unsupervised Learning

1

2

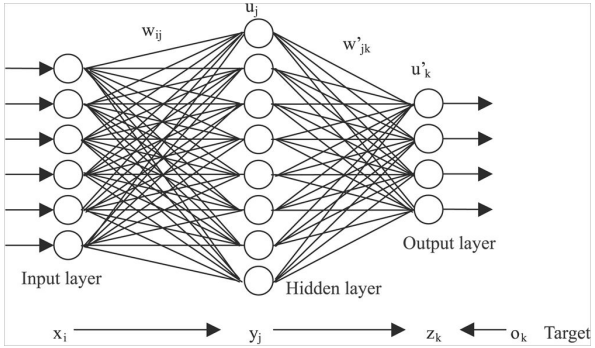

dimensionality reduction & clustering

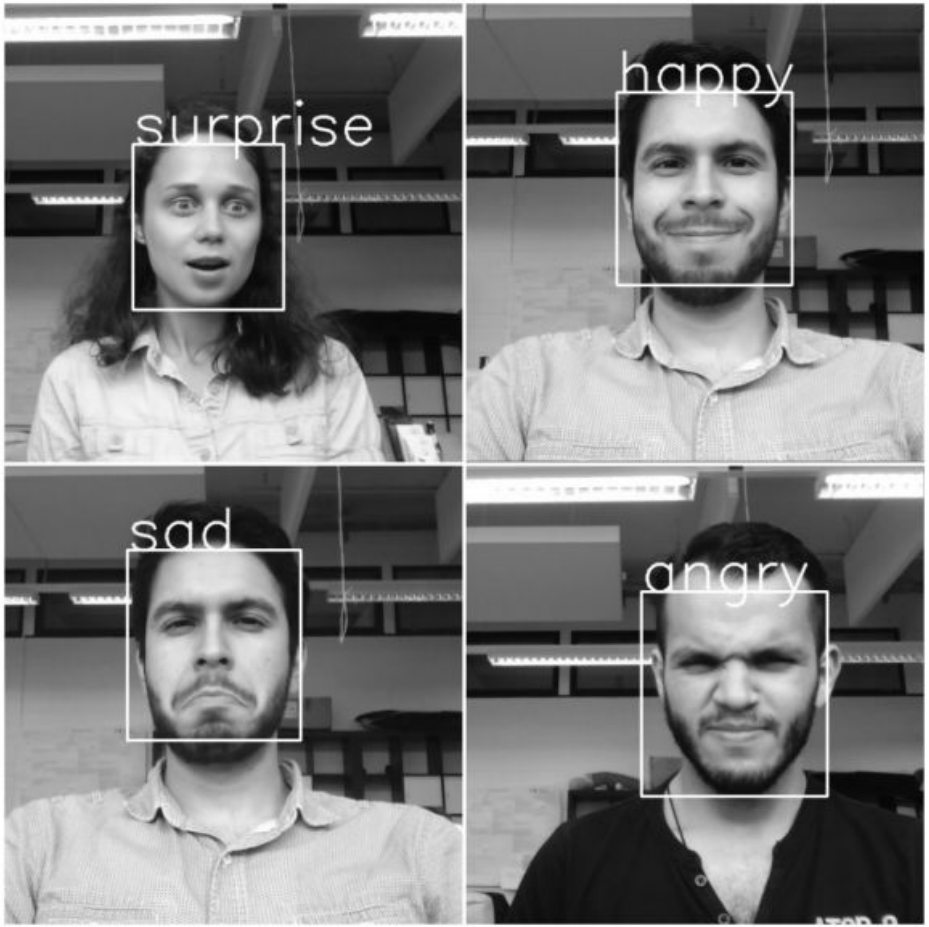

https://github.com/oarriaga/face_classification

Prediction Approach

You want to predict someone's emotion based on an image.

How would you approach this with machine learning?

A
Supervised, Regression

B
Supervised, Classification

C
Unsupervised, dimensionality reduction

D
Unsupervised, clustering

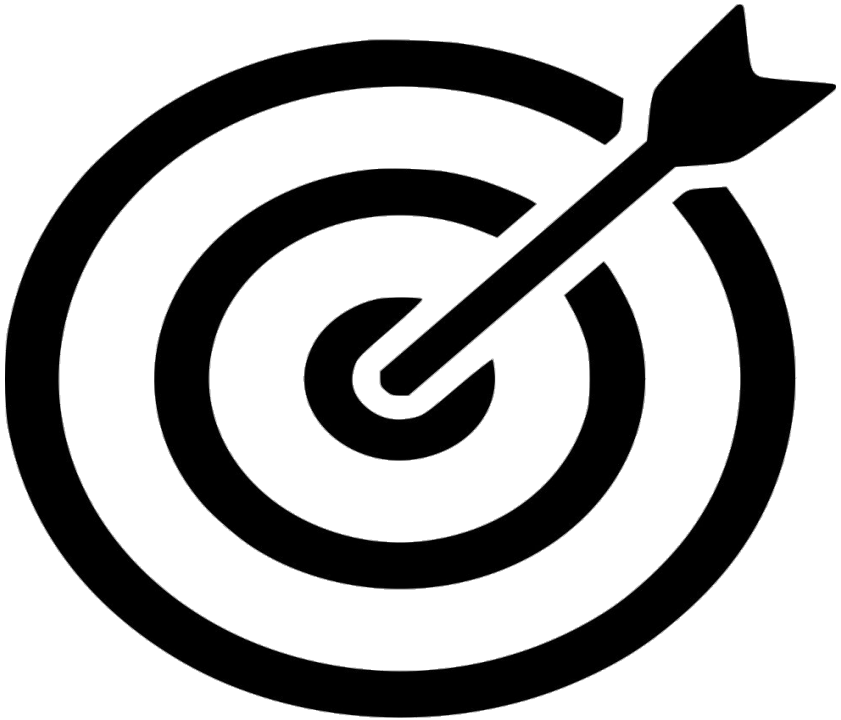

model assessment

continuous variable prediction

Root Mean Squared Error (RMSE)

$$RMSE = \sqrt{\frac{\sum_{i=1}^{N} (Predicted_i - Actual_i)^2}{N}}$$

A few outliers can lead to a big increase in RMSE, even if all the other predictions are pretty good

categorical variable prediction

continuous variable prediction

$$\text{Accuracy} = \frac{\#\text{ of samples predicted correctly}}{\#\text{ of samples predicted}} * 100$$

categorical variable prediction

		Actual	
		Positive	Negative
Predicted	Positive	True Positive (TP)	False Positive (FP)
	Negative	False Negative (FN)	True Negative (TN)

A 2x2 table is a type of confusion matrix

Sensitivity

$$\frac{TP}{TP + FN}$$

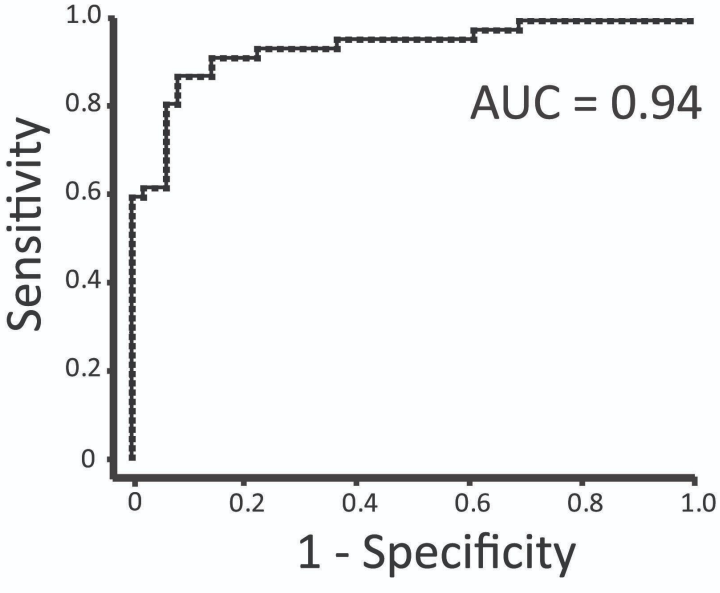

Specificity

$$\frac{TN}{TN + FP}$$

Accuracy	What % were predicted correctly?
Sensitivity	Of those that *were* positives, what % were predicted to be positive?
Specificity	Of those that were *negatives*, what % were predicted to be negative?

Prediction Approach

You've been given a dataset with a number of features and have been asked to predict each individual's age.

What prediction approach would you use?

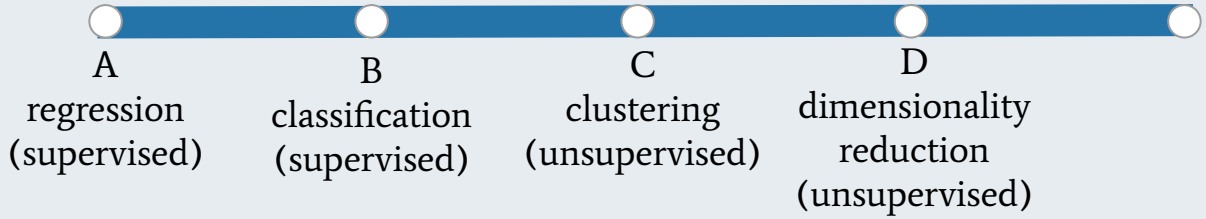

Prediction Approach

After predicting each person's age, how would you assess your model?

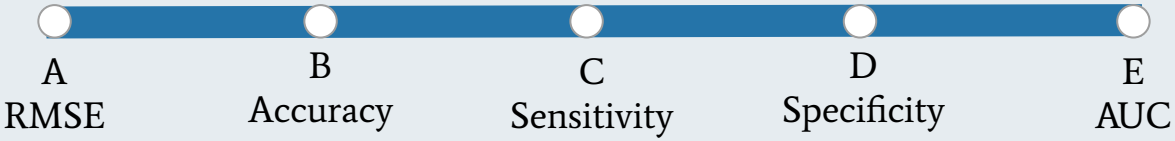

Prediction Approach

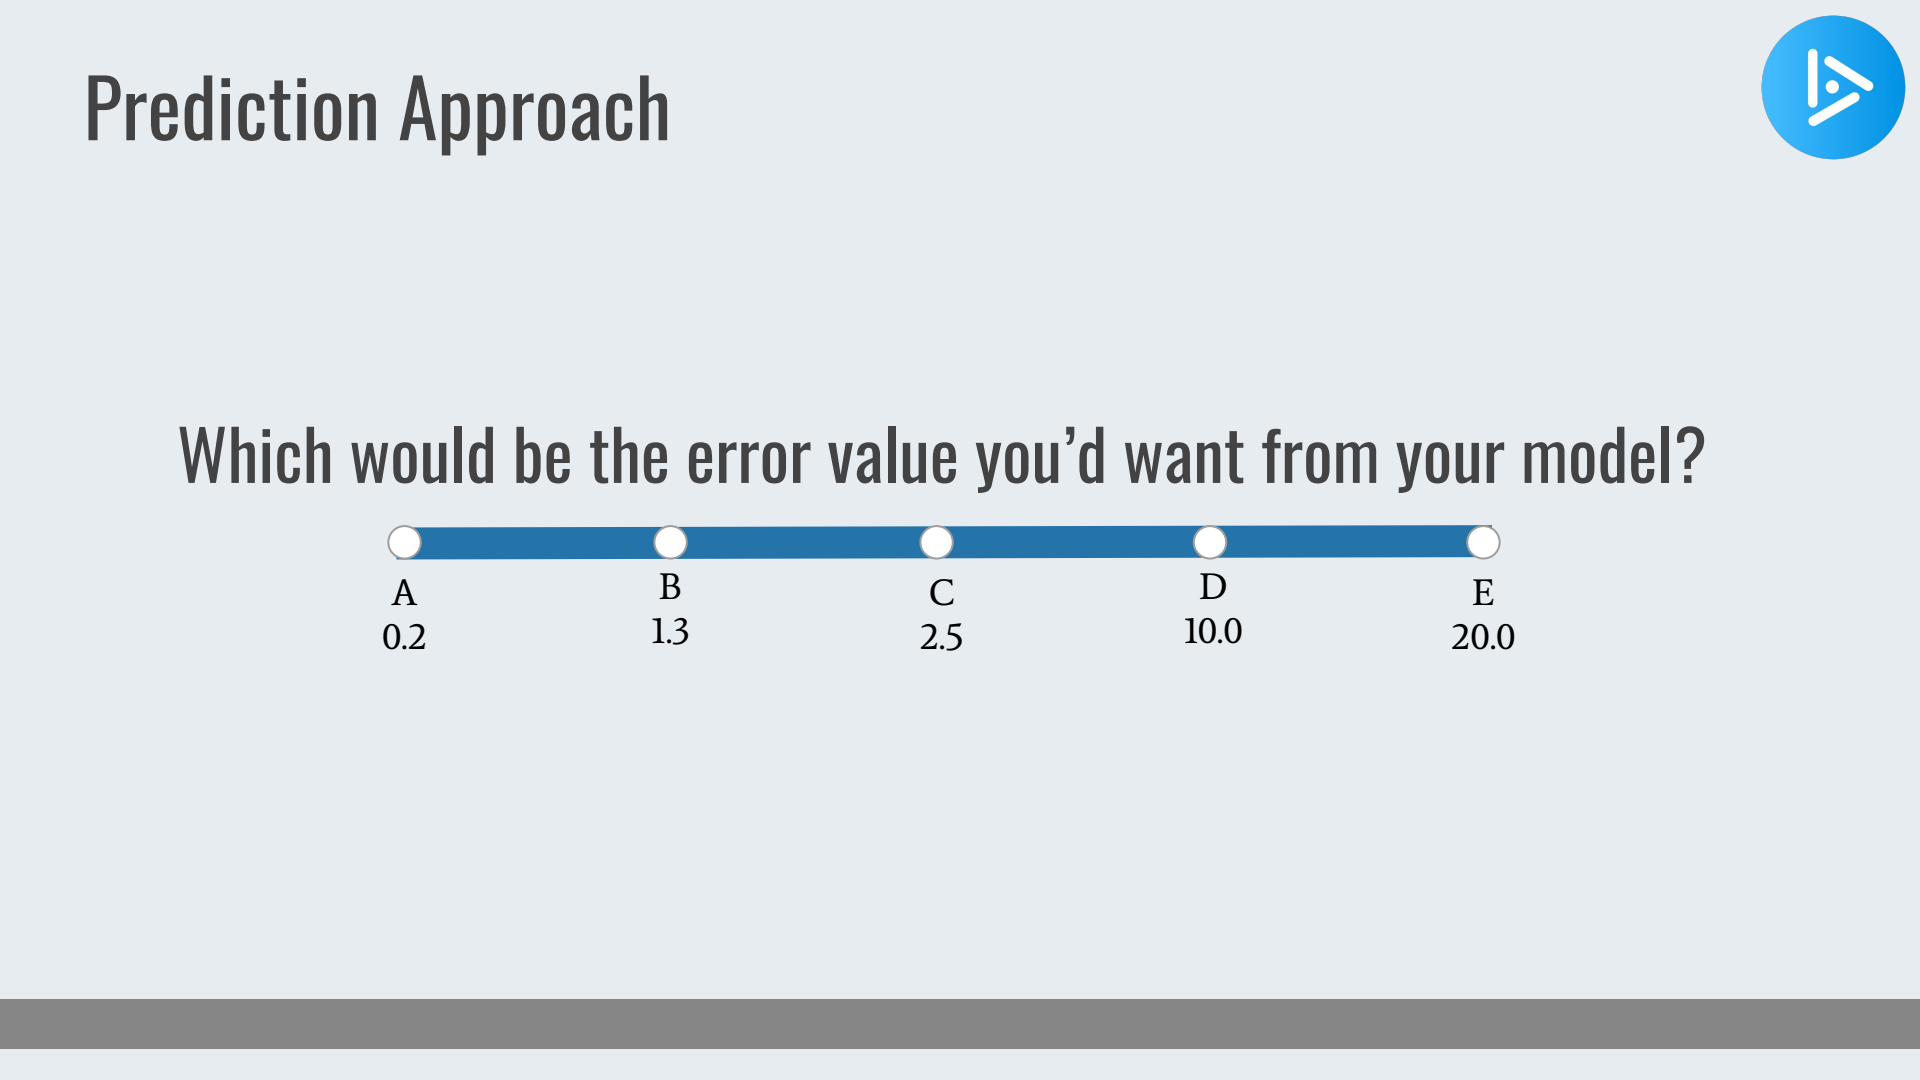

When models are trained on historical data, predictions will perpetuate historical biases

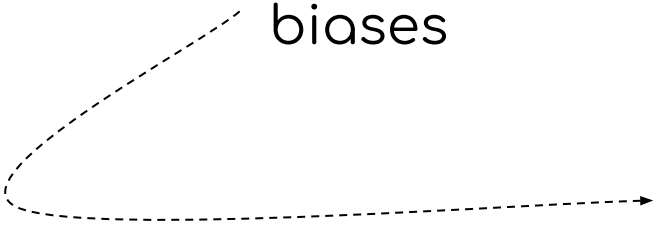

Predictive Analysis Ethics

Dare Obasanjo
@Carnage4Life

Product leader at Microsoft. My team is responsible for advertiser experience for Bing Ads; mobile apps, web UX, desktop apps & SDKs.

Dare Obasanjo
@Carnage4Life

Follow

Machine learning algorithms are driven more by the training data than math. Give an algorithm biased data then results will be biased. E.g.
• Amazon's resumé referral algo which auto rejected women
• Search ads algo which showed background check ads for "black sounding names"

Ryan Saavedra @RealSaavedra
Socialist Rep. Alexandria Ocasio-Cortez (D-NY) claims that algorithms, which are driven by math, are racist

8:59 PM - 22 Jan 2019

https://twitter.com/carnage4life/status/1087892442349678592?s=21

BUSINESS NEWS OCTOBER 9, 2018 / 8:12 PM / A MONTH AGO

Amazon scraps secret AI recruiting tool that showed bias against women

Jeffrey Dastin

8 MIN READ

SAN FRANCISCO (Reuters) - Amazon.com Inc's (AMZN.O) machine-learning specialists uncovered a big problem: their new recruiting engine did not like women.

https://www.reuters.com/article/us-amazon-com-jobs-automation-insight/amazon-scraps-secret-ai-recruiting-tool-that-showed-bias-against-women-idUSKCN1MK08G

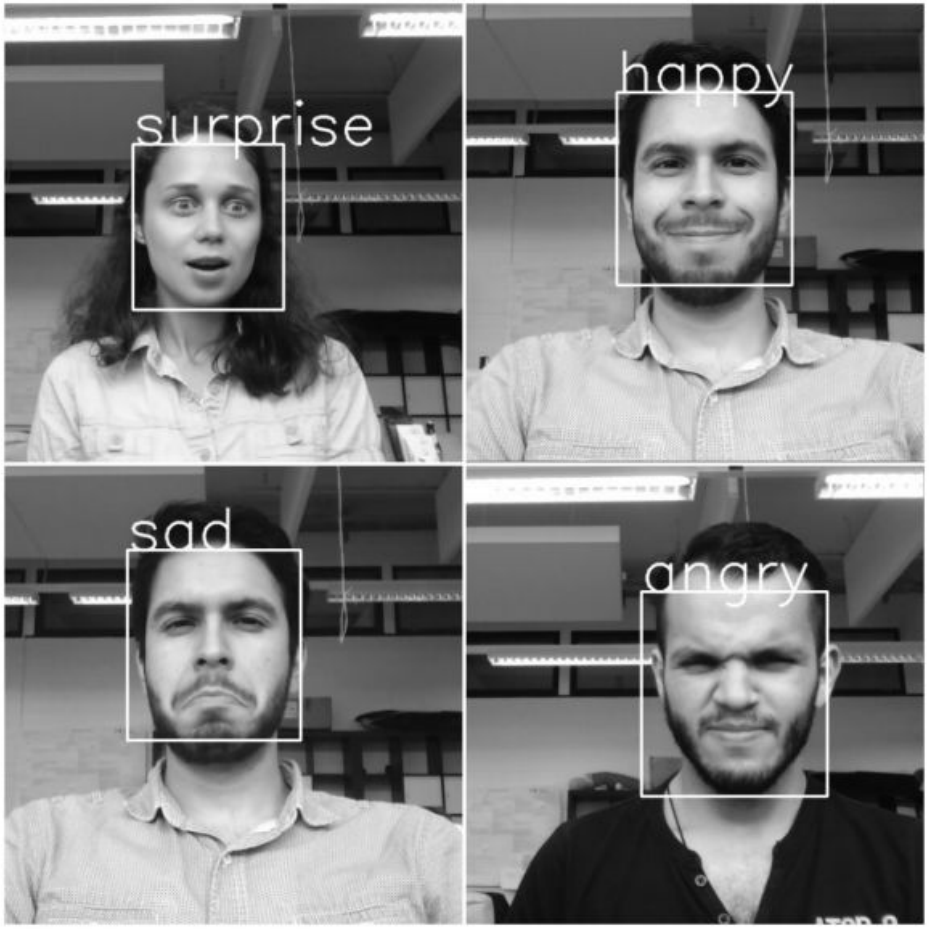

https://github.com/oarriaga/face_classification

https://twitter.com/nke_ise/status/897756900753891328

What to do about bias...

1. Anticipate and plan for potential biases before model generation. Check for bias after.
2. Have diverse teams.
3. Use machine learning to improve lives rather than for punitive purposes.
4. Revisit your models. Update your algorithms.
5. You are responsible for the models you put out into the world, unintended consequences and all.

Discussed so far…

- data partitioning
- feature selection
- supervised & unsupervised machine learning
  - Continuous variables: regression (supervised) and dimensionality reduction (unsuperfied)
  - Categorical variables: classification (supervised; decision trees) or clustering (unsupervised)
- model assessment
  - Continuous: RMSE (& Accuracy)
  - Categorical: Accuracy, Sensitivity, Specificity, AUC
- biased data can & will lead to biased predictions

Data Science Question

Based on data I have about you all, can I predict who in this course will be successful?

Prediction Approach

Which would be the most predictive of your future success?

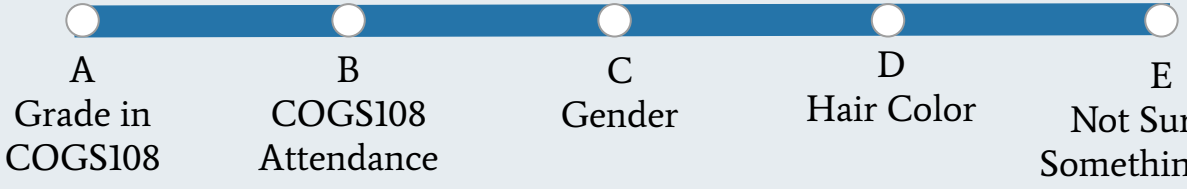

N = 254

train the model:
N = 178
(70% of the data)

test the model:
N = 76
(30% of the data)

train the model

predicted success in test set

	Accuracy	Sensitivity	Specificity
training set	71.2%	76%	67%
test set	49.1%	40%	60%

Assess Prediction Model

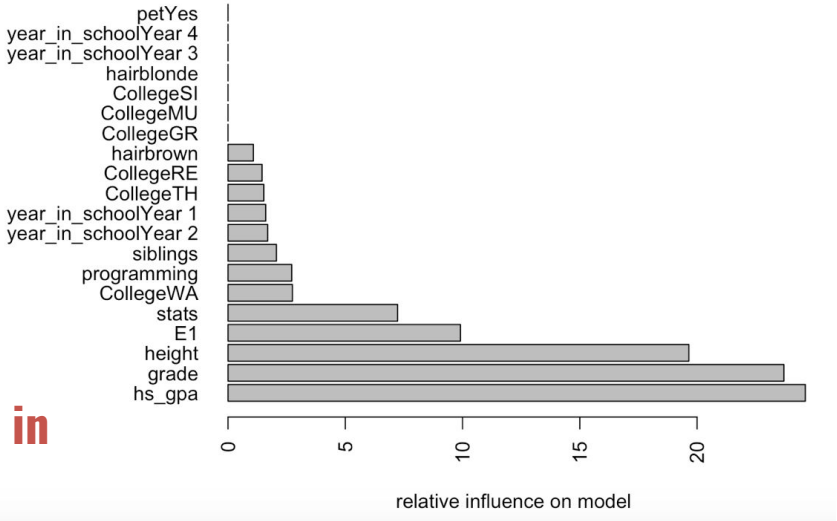

N = 254

train the model:
N = 178
(70% of the data)

test the model:
N = 76
(30% of the data)

train the model

predicted success in test set

	Accuracy	Sensitivity	Specificity
training set	100%	100%	100%
test set	100%	100%	100%

Assess Prediction Model

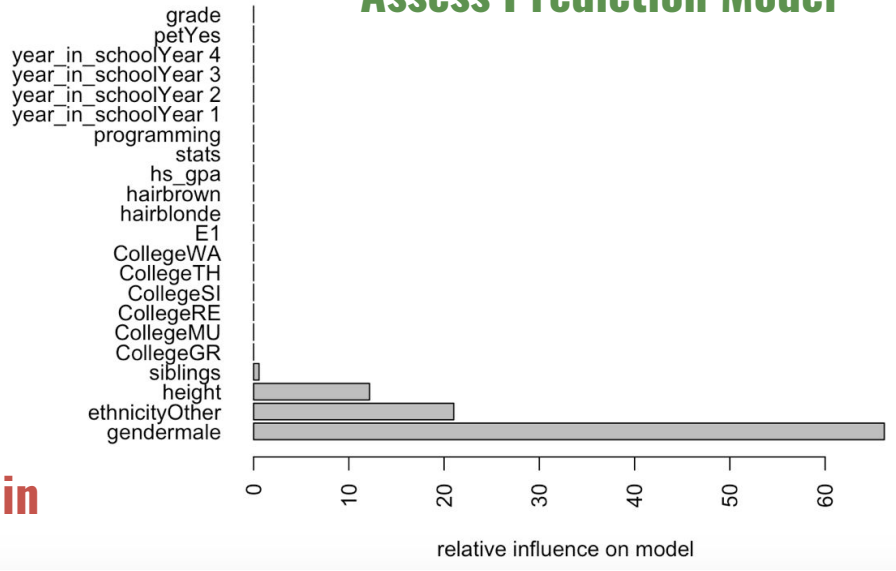

What if I were using these data to determine who I should write recommendation letters for?

Or to determine which students I focus my attention on?

Or whose projects I read?

Or who I allow to come to office hours?

Or who UCSD allows to be data science majors?

Think about whether the models you're building should even be built.